

UNIVERSITÀ
DEGLI STUDI
DI PADOVA



MASTER THESIS IN COMPUTER ENGINEERING

Open Data for Italian Municipalities: Ontology, Data and WebApps

MASTER CANDIDATE

Luca Martinelli

Student ID 2005837

SUPERVISOR

Prof. Gianmaria Silvello

University of Padova

ACADEMIC YEAR
2021/2022

*To my parents and friends,
who have supported (and put up with) me
all these years*

Abstract

Linked Open Data will help transform eGovernment by enabling public administrations to define relationships between data of others. In addition, publishing Linked Open Data creates new knowledge and encourages creativity and innovation. Indeed, governments can activate smarter and more efficient public services and applications, and organizations and citizens can develop new applications and tools in order to work with, analyze, and make sense of the data. In this context, this thesis will present the design and the development of an ontology to model the data produced by municipalities. In addition, the thesis will present the RDF Graph Builder, a software that aims to simplify the publication of resources in the form of Linked Open Data, and *Data Reports*, a web application to for viewing and consulting Linked Open Data through charts, tables and maps.

Sommario

I Linked Open Data contribuiranno a trasformare l'eGovernment, consentendo alle pubbliche amministrazioni di collegare i propri dati ai dati di altri soggetti. Inoltre, la pubblicazione di Linked Open Data crea nuova conoscenza e incoraggia la creatività e l'innovazione. In questo modo, infatti, i governi possono attivare servizi più efficienti e intelligenti e le organizzazioni e i cittadini possono sviluppare nuove applicazioni e strumenti per lavorare, analizzare e dare un senso ai dati pubblici. In questo contesto, questa tesi presenterà la progettazione e lo sviluppo di un'ontologia per modellare i dati prodotti dai comuni. Inoltre, la tesi presenterà l'RDF Graph Builder, un software che mira a semplificare la pubblicazione di risorse sotto forma di Linked Open Data, e *Data Reports*, un'applicazione web per la visualizzazione e la consultazione dei Linked Open Data attraverso grafici, tavole e mappe.

Contents

List of Figures	xi
List of Tables	xiii
List of Code Snippets	xv
List of Acronyms	xvii
1 Introduction	1
1.1 Scope and organization of the thesis	2
2 Background	5
2.1 The Web of Data	5
2.2 The five stars of Open Data	7
2.3 RDF, OWL, and serialization formats	8
2.4 SPARQL	14
2.5 Protégé	15
2.6 Virtuoso	16
2.7 CKAN	20
2.8 OntoPiA	21
3 Related works	25
3.1 Italian cities	25
3.2 European and global cities	27
4 Requirements analysis	31
5 Description of the OntoIM Ontology	33
5.1 Overall design principles	33
5.1.1 Semantic areas	35

CONTENTS

5.1.2	Controlled vocabularies	36
5.2	Area-by-Area	36
5.2.1	Demographic Observations and Events	37
5.2.2	Facilities and Cadastral Data	38
5.2.3	Organizations and Associations	39
5.2.4	Transparency	41
5.2.5	Roads and Traffic	41
5.2.6	Schools	44
5.2.7	Green Zones and Plants	45
5.2.8	Hospitals	46
5.2.9	Waste Production	46
6	RDF Graph Builder	55
6.1	OntoPiA-Py and OntoIM-Py libraries	56
6.2	Data mapping for different semantic areas	64
6.2.1	Addresses	66
6.2.2	Organizations	69
6.2.3	Schools	71
7	Web Applications	75
7.1	CKAN	75
7.2	Data Reports	78
8	Conclusions and Future Works	85
References		87
Acknowledgments		89

List of Figures

2.1	The structure of a triple, with two nodes and a predicate connecting them.	9
2.2	The example of the Resource Description Framework (RDF) Graph presented by World Wide Web Consortium (W3C).	10
2.3	A snapshot of the Protégé "Active ontology" tab.	16
2.4	A snapshot of the Protégé "Entities" tab.	17
2.5	A snapshot of the Virtuoso SPARQL Protocol and RDF Query Language (SPARQL) endpoint.	18
2.6	A snapshot of the "Quad Store Upload" tab.	18
2.7	Examples of CKAN Open Data governments portals.	20
2.8	The OntoPiA ontological stack.	22
5.1	Demographic Observations and Events semantic area.	47
5.2	Facilities and Cadastral Data semantic area.	48
5.3	Organizations and Associations semantic area.	48
5.4	Transparency semantic area.	49
5.5	Roads and Traffic semantic area.	50
5.6	Schools semantic area.	51
5.7	Green Zones and Plants semantic area.	52
5.8	Hospitals semantic area.	52
5.9	Waste Production semantic area.	53
6.1	RDF Graph Builder architecture.	56
7.1	A snapshot of the Comune di Sona's CKAN Open Data portal. . . .	76
7.2	A snapshot of the Data Reports home page.	79
7.3	Snapshots for different Data Reports's pages.	81

List of Tables

2.1	The main modeling constructs provided by RDF Schema.	11
2.2	A query result example from DBpedia.	14
2.3	Ontologies part of the OntoPiA network.	24
3.1	Analysis of Italian cities' Open Data Portals. The data reported in this table was collected during April 2022.	26
3.2	Analysis of European and Global cities' Open Data Portals. The data reported in this table was collected during April 2022.	28
5.1	The data collected as reference for designing the OntoIM ontology, and their source.	34

List of Code Snippets

6.1	Example of a config.ini file that defines sources for some semantic areas.	55
6.2	Part of the ns.py file that contains the namespaces of OntoPiA's ontologies.	57
6.3	Part of the __init__.py file that contains the two functions for creating and saving the graph.	58
6.4	The ontopia-py's Thing class.	60
6.5	The ontopia-py's Sequence and Collection classes. Thanks to the object-oriented programming inheritance it is possible to map ontology classes and properties into Python classes and attributes.	61
6.6	The ontopim-py's Organization class. This class inherits the one defined in ontopia-py, declaring the new attributes.	61
6.7	An example of the creation of an RDF Graph with the ontoim-py and ontopia-py libraries.	63
6.8	The common functions in the utils package.	65
6.9	The part of the RDF Graph Builder that inserts the streets toponyms into the graph, and the config.ini file relative to the addresses.	68
6.10	The function that retrieve the street identifiers from the string of the address.	70
6.11	The part of the code that build the RDF Graph for demographic observations on the school.	72
7.1	Configuration template file for the CKAN Open Data portal.	77
7.2	"Associazioni" report page in Markdown.	80
7.3	The querySPARQL function.	81
7.4	The sunburst chart about the number of citizens by locality and gender.	82
7.5	The DataReports configuration file.	84

List of Acronyms

AgID Agency for Digital Italy

API Application Programming Interface

ASCII American Standard Code for Information Interchange

CSS Cascade StyleSheet

CSV Comma Separated Values

DBMS DataBase Management System

HTML HyperText Markup Language

HTTP HyperText Transfer Protocol

IRI International Resource Identifier

JSON JavaScript Object Notation

LOD Linked Open Data

OntoIM Ontology for Italian Municipalities

OWL Web Ontology Language

RDF Resource Description Framework

RDFa RDF in Attributes

RDFS RDF Schema

SEO Search Engine Optimization

SPARQL SPARQL Protocol and RDF Query Language

LIST OF CODE SNIPPETS

SQL Structured Query Language

TSV Tab Separated Values

URI Uniform Resource Identifier

URL Uniform Resource Locator

W3C World Wide Web Consortium

XML eXtensible Markup Language

1

Introduction

Linked Open Data will help transform eGovernment by enabling public administrations to define relationships between data of others. In addition, publishing Linked Open Data creates new knowledge and encourages creativity and innovation. Indeed, governments can activate smarter and more efficient public services and applications, and organizations and citizens can develop new applications and tools in order to work with, analyze, and make sense of the data [ISA13].

In this context, the Ontology for Italian Municipalities (OntoIM) has the aim to extend the Italian OntoPiA ontology in order to describe the main semantic areas of a city, enabling it to publish the data as Linked Open Data. Such areas are: Accommodation Facilities, Addresses and Civic Numbers, Cultural Heritage, Demographic Observations and Events, Facilities and Cadastral Data, Green Zones and Plants, Hospitals, IoT Events, Organizations and Associations, Point of Interests, Public Events, Public Services, Roads and Traffic, Schools, Transparency, Projects, and Public Contracts, Waste Production.

The OntoIM ontology will also allow to link the data with other resources on the web, from national agencies, and from other cities.

In addition to the OntoIM ontology, this thesis has the purpose to provide tools and web applications with the purpose of simplifying on the one hand the publication of Open Data and Linked Open Data by public administrations and municipalities, and on the other hand making those data accessible through the use of data visualization tools.

The project would have benefits: (1) **for citizens**, who will be able to more

1.1. SCOPE AND ORGANIZATION OF THE THESIS

easily access the data produced by the public administration, better understand their local area, and obtain information on events or accommodations; (2) **for businesses**, which will have data on the area that they can analyze, for example, to make targeted marketing plans; and (3) **for the public administration**, which would have data enriched with semantic meaning through the use of ontology, have the opportunity to be more transparent, could use the data to make and plan more targeted projects, and provide smarter and more efficient public services and applications.

Finally, a use case with data from the Comune di Sona¹ will also be shown in this thesis. In fact, this project will be part of the *Innovation Lab*² project, a project funded by the Regione Veneto that aims to spread digital and Open Data culture.

1.1 SCOPE AND ORGANIZATION OF THE THESIS

The goal of this thesis is to design and develop:

1. The first version of the OntoIM ontology;
2. Two Python libraries, `ontoim-py` and `ontopia-py` with the aim of simplifying the creation of the RDF Graph;
3. A CKAN Open Data portal for publishing data on the web;
4. *Data Reports*, a web application to publish reports with charts, tables, and maps using data obtained from the SPARQL endpoint.

The following chapters will focus on the various steps that were required to achieve these goals, in detail:

Chapter 2: Background This chapter will present an introduction on the Semantic Web, Linked Open Data, and its technologies such as RDF, OWL, serialization formats, and SPARQL. This chapter will also introduce the tools used for developing the ontology, and managing and publishing the resources: Protégé, Virtuoso, and CKAN. Finally, the chapter will present OntoPiA, the Italian ontology that the OntoIM ontology imports and extends.

¹<https://comune.sona.vr.it>

²<https://innovationlab.regionev.it/>

Chapter 3 This chapter presents the results of an analysis conducted on Italian, European and global cities regarding their approach to Open Data and in particular Linked Open Data.

Chapter 4 This chapter will explain the ontology design approach, based on the analysis of data collected, from the analysis on data published by Italian cities, and from the needs and requirements of the Comune di Sona.

Chapter 5 This chapter will elaborate on the approach followed to design the ontology, and will describe in details the main classes and properties contained in each of the semantic areas that compose the OntoIM ontology.

Chapter 6 This chapter covers the development of the Python libraries `ontoim-py` and `ontopia-py`, and the use of these libraries to create the RDF Graph using examples with data collected and published by the Comune di Sona.

Chapter 7 This chapter will present the design and the development of the CKAN Open Data portal, and the *Data Reports* web application.

Chapter 8 This chapter presents the overall conclusions of the thesis, the results achieved, and the future developments of the project.

2

Background

2.1 THE WEB OF DATA

The World Wide Web was originally designed to be a space where documents are connected by links without semantic value, and most of these documents are designed for humans to read, not for machines to process. For this reason, Tim Berners-Lee in 2001 introduced the idea of the Semantic Web. In particular, the Semantic Web is an enhancement of the current Web that aims to create a web of data, in which information has a well-defined meaning and can be easily read and processed by programs [BHL01].

Due to this machine-comprehensible capacity, the Semantic Web has enormous potential to automate daily tasks in our lives and is helping to advance scientific and health care fields [Fei+07], such as drug discovery and clinical research, but also in the automotive industry, in the enhancement of cultural heritage, etc.¹ In this context, ontologies play a fundamental role in supporting interoperability and common understanding between different web applications and services, solving the problem of semantic heterogeneity [Tay10].

Although there are different definitions of "ontology" [Tay10], in computer science, an ontology is defined as an "explicit and formal specification of a shared conceptualization" [Gru95], where conceptualization means a simplified view of

¹<https://www.w3.org/2001/sw/sweo/public/UseCases/>

2.1. THE WEB OF DATA

the world we wish to represent. An ontology is made up of four main types of components, which are (1) *classes* (or *concepts*), which describe concepts in the domain; (2) *instances* of classes, which represent specific objects or elements of a class; (3) *properties* (or *slots*), which are used to express relationships between a first concept in the domain and a second concept in the range; (4) *axioms* (or *role restrictions*), which are used to impose constraints on the values of instances and classes [Tay10; NM+01]. In addition to the interoperability problem, ontologies are also used to satisfy the following needs:

- To share common understanding of the structure of information among people or software agents;
- To enable reuse of domain knowledge;
- To make domain assumptions explicit;
- To separate domain knowledge from the operational knowledge
- To analyze domain knowledge [NM+01].

Along with ontologies, controlled vocabularies, taxonomies, and thesauri are other resources used in different domains, in particular the medical one [IB14]. A controlled vocabulary is a closed list of named subjects, called *terms*, which is usually used for classification. A taxonomy is a subject-based classification that organizes terms in a controlled vocabulary into a hierarchy. Finally, a thesaurus extends taxonomies, allowing making other statements about the subjects and providing a much richer vocabulary [IB14].

In 2006, Tim Berners-Lee used for the first time the term Linked Data to describe the structured and interlinked data that populate the Semantic Web [Ber06]. He also introduced a set of rules, also known as the Linked Data Principles, to provide some best practices for publishing and connecting data on the Web [BHB11]. These principles, published by W3C, are the following:

1. Use URIs as names for things;
2. Use HTTP URIs so that people can look up those names;
3. When someone looks up URIs, provide useful information, using standards such as RDF and SPARQL;
4. Include links to other URIs so that they can discover more things.

The two main fundamental technologies for Linked Data are Uniform Resource Identifiers (URIs) and HyperText Transfer Protocol (HTTP). In particular, URIs are used to identify any entity that exists in the world, while HTTP provides a simple and universal mechanism for retrieving the resources to which they refer. These two technologies are integrated in RDF, which provides a graph-based data model to structure and link data that describe entities in the world [BHB11]. Using HTTPs, URIs, and RDF, Linked Data builds on the architecture of the Web, called the Web of Data. This means that the Web of Data shares many properties with the traditional Web, which are:

- Web of Data can contain any type of data;
- Anyone can publish data on the Web of Data;
- Publishers are not restricted in the choice of vocabularies used to represent the data;
- Entities are connected by RDF links [BHB11].

However, in addition to those of the traditional Web, the Web of Data also has the following characteristics:

- Data are separated from formatting and presentational aspects;
- Data is self-describing;
- Data access is simplified by the use of the HTTP and RDF standards;
- Web of Data is open, and new data sources can be discovered at run-time by following RDF links [BHB11].

Semantic Web and Linked Data are empowered by technologies developed by the World Wide Web Consortium such as RDF, OWL, serialization formats (Section 2.3), and SPARQL (Section 2.4).²

2.2 THE FIVE STARS OF OPEN DATA

Linked Data does not have to be open and can be used internally, such as for personal data. When Linked Data is released under an open license that does not

²https://www.w3.org/2001/sw/wiki/Main_Page

2.3. RDF, OWL, AND SERIALIZATION FORMATS

impede its reuse for free, such as Creative Commons CC-BY³ or the Italian Open Data License⁴, we can use the term Linked Open Data (LOD) [Ber06]. In 2010 Tim Berners-Lee developed a star rating system to define and classify Linked Open Data, "in order to encourage people, especially government data owners, along the road to good linked data" [Ber06]. The star rating system assigns a star if the information is publicly available under an open license, even if the information is a photo or an image scan of a table. The more stars the information gets, the easier it will be for people (and machines) to use it [Ber06].

- ★ Available on the Web (any format) but with an open license to be Open Data
- ★★ Available as machine-readable structured data (e.g., Excel instead of an image scan of a table)
- ★★★ Available in a non-proprietary format (e.g., CSV instead of Excel)
- ★★★★ Use URIs to identify things, so that people can point at your stuff
- ★★★★★ Data are linked to other people's data to provide context

However, as the information receives a greater number of stars, both the benefits for consumers and the costs for the publisher increase. In particular, a five-stars data let consumers discover new data of interest, access to the data schema, reuse parts of the data, and link it to other places. They also do not have to pay for tools in order to read the data (e.g., Excel), and they can download and export the data into other formats and process them. On the other hand, to make these data available, publishers must invest time and resources in slicing and organizing the data, assigning URIs to the data items, thinking about how to represent them, linking the data with other data on the Web and making them discoverable [BK11].

2.3 RDF, OWL, AND SERIALIZATION FORMATS

The Resource Description Framework (RDF) is a W3C graph-based standard model to represent information about resources on the Web (including documents,

³<https://creativecommons.org/>

⁴<https://www.dat.gov/content/italian-open-data-license-v20>

people, physical objects, and abstract concepts). Using RDF, machines can process information on the Web using common parsers and processing tools, and information can be exchanged between different applications without losing meaning [Con+14b]. In particular, in recent years RDF has become the *de-facto* standard for publishing Linked Data on the Web. The core structure of the RDF syntax is a set of statements, called *triples*, because they consist of three elements: a *subject*, a *predicate*, and an *object*, following the structure <subject> <predicate> <object>, which can be visually represented in Figure 2.1 [Con+14a].

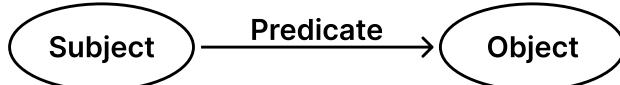


Figure 2.1: The structure of a triple, with two nodes and a predicate connecting them.

The subject and the object represent the two resources being related. The relationship that goes from the subject to the object is called *property*, and its nature is represented by the predicate. A set of statements generate a direct graph, called RDF Graph, where subjects and objects are the nodes of the graph, and the predicates form the arcs. For example, the set of triples below produces the graph shown in Figure 2.2 [Con+14b].

```

<Bob> <is a> <person>.
<Bob> <is a friend of> <Alice>.
<Bob> <is born on> <the 4th of July 1990>.
<Bob> <is interested in> <the Mona Lisa>.
<the Mona Lisa> <was created by> <Leonardo da Vinci>.
<the video 'La Joconde à Washington'> <is about> <the Mona Lisa>
  
```

In an RDF Graph, resources may be represented using an International Resource Identifier (IRI), a *literal value* or a *blank node*. An IRI is a generalization of URI, where non-ASCII characters are allowed in the IRI character string. IRIs identify resources, and can appear in all three positions of a triple. In the example above, the IRI for Leonardo Da Vinci in DBpedia⁵ is http://dbpedia.org/resource/Leonardo_da_Vinci.

⁵<https://www.dbpedia.org/>

2.3. RDF, OWL, AND SERIALIZATION FORMATS

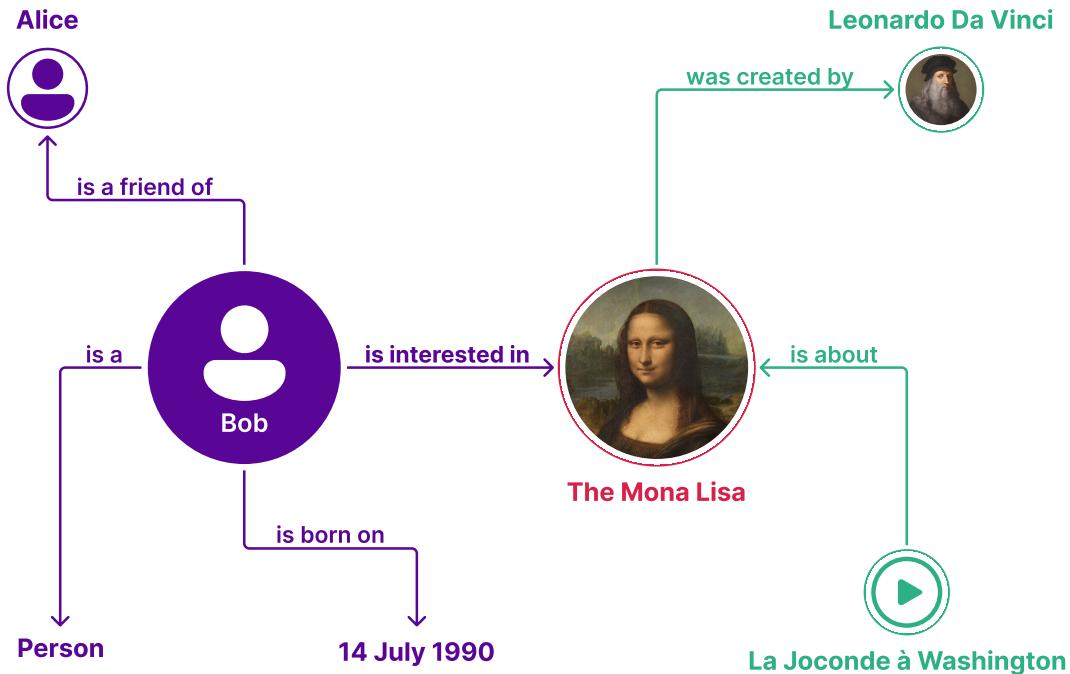


Figure 2.2: The example of the RDF Graph presented by W3C.

Literals are basic values such as strings, dates, and numbers. In the RDF Graph literals can only be used as objects, and consists of two or three elements, which are: (1) the value itself; (2) an IRI that identifies the *datatype* (string, number, date, etc); (3) if and only if the datatype is a `rdf:langString`,⁶ a *language tag* (such as en, it, fr, etc) [Con+14a].

Finally, blank nodes can appear in the subject and object position of a triple and are used to represent resources without using a IRI [Con+14b].

In Section 2.1 ontologies and vocabularies are presented as a core element for creating the Semantic Web. The RDF data model does not provide semantic information about the resources. For this reason, RDF provides the RDF Schema (RDFS) language, that allows to define semantic characteristics of data. RDF Schema uses the notion of *class* to classify resources, while uses the *type* property to define a relation between an instance and its class. RDF Schema also allows defining type restrictions on subject and objects of particular triples through *domain* and *range* restrictions. Finally, with RDF Schema it is also possible to define hierarchies of classes and properties, using *subClassOf* and *subPropertyOf*

⁶<http://www.w3.org/1999/02/22-rdf-syntax-ns#langString>

predicates [Con+14b]. All of these modeling constructs provided by RDF Schema are summarized in Table 2.1.

Construct	Syntactic form	Description
Class	C rdf:type rdfs:Class	C is an RDF class
Property	P rdf:type rdf:Property	P is an RDF property
type	I rdf:type C	I is an instance of C
subClassOf	C1 rdfs:subClassOf C2	C1 is a subclass of C2
subPropertyOf	P1 rdfs:subPropertyOf P2	P1 is a sub-property of P2
domain	P rdfs:domain C	domain of P is C
range	P rdfs:range C	range of P is C

Table 2.1: The main modeling constructs provided by RDF Schema.

However, in 2004 the World Wide Web Consortium presented Web Ontology Language (OWL), a more complete language for publishing and sharing ontologies on the Web [Bec+04], and replaced in 2009 and then in 2012 by OWL 2. OWL 2 is a Semantic Web language to represent rich and complex knowledge about things, groups of things, and relations between things. In addition, since OWL is a computational logic-based language, the knowledge expressed in OWL can be reasoned with by computer programs either to verify the consistency of that knowledge or to make implicit knowledge explicit. A OWL document, called *ontology*, can be published in the World Wide Web and may refer to or be referred from other OWL ontologies [Hit+09]. In OWL 2 knowledge is represented by statements, called *axioms*. Axioms normally refer to objects of the world and describe them by putting them into categories or saying something about their relation. In OWL 2 objects, categories and relations are called *entities*, and in particular objects are denoted as *individuals*, categories as *classes* and relations as *properties*. Moreover, properties are further subdivided into (1) *object properties* that relate objects to objects; (2) *datatype properties* that assign data values to objects; (3) *annotation properties* that encode information about the ontology itself. Finally, names of entities can be combined into *expressions* using *constructors* to form complex descriptions from basic ones [Hit+09].

2.3. RDF, OWL, AND SERIALIZATION FORMATS

In order to publish RDF data on the Web, the RDF Graphs need to be serialized. Today there are several serialization formats, but the most famous one are: N-Triples, Turtle, RDF/XML, RDFa, and JSON-LD. These formats are briefly described below, reporting as example of small excerpt of DBpedia⁷ is reported.

N-Triples⁸ It's one of the simplest formats, formed by sequences of RDF triples. Each statement is formed by the subject, predicate, object, and a ".", that are separated by white space.

```
<http://dbpedia.org/page/Jotaro_Kujo>
<http://dbpedia.org/ontology/relative>
<http://dbpedia.org/page/Joseph_Joestar> .
```

Turtle⁹ It's a common data format for serializing RDF Graphs that introduces some features to N-Triples language. In particular, it introduces the use of @base IRI and relative IRIs, @prefix and prefixed names, predicate lists separated by ";", object lists separated by ",", and the representation of rdfs:type with the token a.

```
@prefix dbr: <http://dbpedia.org/page/> .
@prefix dbo: <http://dbpedia.org/ontology/> .
```

```
dbr:Jotaro_Kujo dbo:relative dbr:Joseph_Joestar .
```

RDF/XML¹⁰ Expresses RDF Graphs as an XML document. The nodes and predicates are represented in XML terms: element names, attribute names, element contents and attribute values.

```
<rdf:RDF xmlns:dbr="http://dbpedia.org/page/"
           xmlns:dbo="http://dbpedia.org/ontology/"
           xmlns:rdf="http://www.w3.org/1999/02/22-rdf-syntax-ns#"
           xml:base="http://www.ldf.fi/service/rdf-serializer/">
  <rdf:Description
    rdf:about="http://dbpedia.org/page/Jotaro_Kujo">
```

⁷<https://www.dbpedia.org/>

⁸<https://www.w3.org/TR/n-triples/>

⁹<https://www.w3.org/TR/turtle/>

¹⁰<https://www.w3.org/TR/rdf-syntax-grammar/>

```

<dbo:relative
  rdf:resource="http://dbpedia.org/page/Joseph_Joestar"/>
</rdf:Description>
</rdf:RDF>
```

RDF in Attributes (RDFa)¹¹ Provides a set of markup attributes to HTML pages to augment the visual information on the Web with machine-readable hints.

```

<body
  prefix="dbr: http://dbpedia.org/page/
  dbo: http://dbpedia.org/ontology/">
<div about="dbr:Jotaro_Kujo">
  <div
    rel="dbo:relative"
    resource="dbr:Joseph_Joestar">
  </div>
</div>
</body>
```

JSON-LD¹² Serializes RDF Graphs into JavaScript Object Notation (JSON). The syntax is designed to easily integrate into deployed systems that already use JSON. It's intended to be a way to use Linked Data in Web-based programming environments, to build interoperable Web services, and to store Linked Data in JSON-based storage engines.

```
[
  {
    "@id": "http://dbpedia.org/page/Joseph_Joestar"
  },
  {
    "@id": "http://dbpedia.org/page/Jotaro_Kujo",
    "http://dbpedia.org/ontology/relative": [
      {
        "@id": "http://dbpedia.org/page/Joseph_Joestar"
```

¹¹<https://www.w3.org/TR/rdfa-primer/>

¹²<https://www.w3.org/TR/json-ld/>

2.4. SPARQL

```
    }
]
}
]
```

2.4 SPARQL

SPARQL is a query language developed by W3C retrieve and manipulate RDF Graph content on the Web or in an RDF store. A SPARQL query contains a set of triple patterns called *basic graph pattern*. These patterns are like RDF triples except that subjects, predicates and objects may be replaced by variables. The basic graph pattern matches a subgraph of the RDF data and returns a new RDF Graph in which the variables are replaced with the matched data. Queries are usually processed by an HTTP service, called *SPARQL endpoint*. [Con+13]. The example below shows a SPARQL query on DBpedia SPARQL endpoint,¹³ while Table 2.2 shows its result.

```
PREFIX dbr: <http://dbpedia.org/page/>
PREFIX dbo: <http://dbpedia.org/ontology/>
PREFIX dbp: <http://dbpedia.org/property/name/>
```

```
SELECT ?relative ?name WHERE {
    dbr:Jotaro_Kujo dbo:relative ?relative .
    ?relative dbp:name ?name .
}
```

relative	name
dbr:Dio_Brando	"Dio Brando"@en
dbr:Joseph_Joestar	"Joseph Joestar"@en
dbr:Jonathan_Joestar	"Jonathan Joestar"@en

Table 2.2: A query result example from DBpedia.

SPARQL queries supports features like union of patterns, nesting queries,

¹³<https://dbpedia.org/sparql>

optional patterns or filtering values. Once the RDF subgraph is computed, it's also possible to modify it by ordering, limiting and grouping the values.

Another important feature of SPARQL is the possibility to perform federated queries, which explicitly delegates certain subqueries to different SPARQL endpoints, allowing to navigate through the Web of Data.

Finally, to return a more machine-readable form, SPARQL supports four common exchange formats, which are: eXtensible Markup Language (XML), JSON, Comma Separated Values (CSV), and Tab Separated Values (TSV) [Con+13].

2.5 PROTÉGÉ

Protégé¹⁴ is the most popular free and open source¹⁵ ontology development environment.¹⁶ The first version was developed by Mark Musen in 1987, and has been so far by a team at Stanford University [Gen+03]. The latest version, 5.5.0, has been released in March 2019, and it is written in Java, making it a cross-platform tool. In recent years, in addition to the desktop version, a web version, called WebProtégé¹⁷ is also being developed, focused on collaborative viewing and editing.

Protégé supports creation and editing of one or more ontologies, providing a customizable graphic user interface. Among the several features available in Protégé, the most relevant are the possibility to create, rename and delete entities, add notations, merge ontologies, and more. It also includes a visualization tool for interactive navigation of ontology relationships and different reasoners.

The two most important sections for creating and editing an ontology are the "Active ontology" and "Entities" tabs. The first, that is opened by default, is designed to view and edit the information of the ontology, such as its IRI, its annotations and the imported ontologies. On the right there is also a panel that reports some metrics about the ontology, such as the total number of axioms, classes, properties, and more. Figure 2.3 shows an example of the "Active ontology" tab.

The "Entities" tab is the most important section for creating an ontology. Indeed, in this tab it is possible to manage the classes, the properties (object proper-

¹⁴<https://protege.stanford.edu/>

¹⁵<https://github.com/protegeproject/protege>

¹⁶<https://protege.stanford.edu/shortcourse/>

¹⁷<https://webprotege.stanford.edu/>

2.6. VIRTUOSO




Figure 2.3: A snapshot of the Protégé "Active ontology" tab.

ties, data properties, and annotation properties), datatype and individuals. The left part provides a navigation tool to select, add and delete entities, while the right part is focused on viewing and editing the selected entity by adding properties and axioms. Figure 2.4 shows an example of the "Entities" tab.

Of course, these were only the most relevant tools of Protégé, whose full documentation is available at <http://protegeproject.github.io/protege/>.

2.6 VIRTUOSO

Virtuoso Universal Server,¹⁸ often called just Virtuoso, or OpenLink Virtuoso, at core is a high-performance object-relational SQL database. It was born in 1998 when OpenLink Software wanted to merge in a single solution its Universal Data Access Middleware and Kubl DBMS.¹⁹

Besides the database, Virtuoso has a built-in web server with support to Virtuoso's Web Language (VSP), and the most popular scripting languages such as PHP or ASP.NET. This same web server provides SOAP and REST access to Virtuoso stored procedures, supporting a broad set of WS* protocols. Virtuoso has also a

¹⁸<https://virtuoso.openlinksw.com/>

¹⁹<https://vos.openlinksw.com/owiki/wiki/VOS/VOSHISTORY>

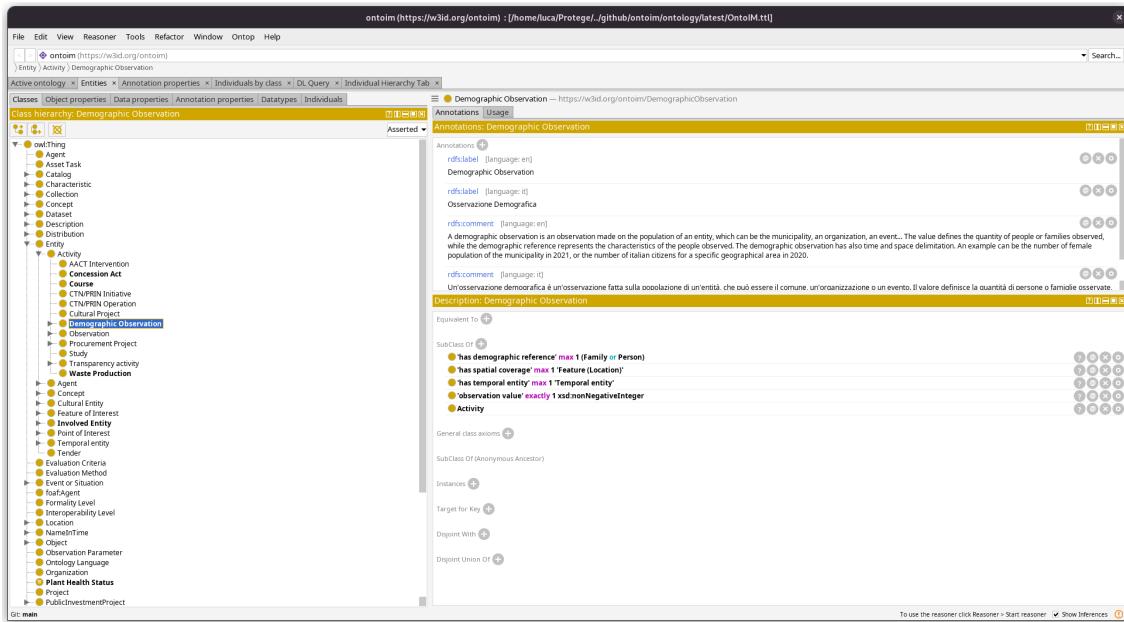


Figure 2.4: A snapshot of the Protégé "Entities" tab.

built-in WebDAV repository to host static and dynamic web content and provide versioning, making it a convenient and secure place for keeping files on the net.²⁰

Since 2005, Virtuoso supports SPARQL for querying RDF data stored in its Quad Store database. In particular, it supports the HTTP-based SPARQL Protocol, SPARQL federated queries, different exchange formats such as HTML, CSV, TSV, JSON, RDF/XML, Turtle, N-Triples, and more. For this reasons Virtuoso has become the most popular and efficient tool for serving a SPARQL endpoint, which is usually located at `http://{host}/sparql`. Figure 2.5 shows an example of how the endpoint looks like.

All the aspects of a Virtuoso instance can be managed through the Virtuoso Conductor, that is located at `http://{host}/conductor`. For example, from "Linked Data" tab it is possible to add and remove RDF Graphs, import schemas, declare persistent namespaces, generate statistics such as the number of classes, triples, subjects, etc.

There are many methods to insert an RDF resource into the Virtuoso Quad Store. Some of them are:

Virtuoso Conductor Using Virtuoso Conductor web interface, under "Linked Data" and then "Quad Store Upload" tab it is possible to upload a RDF

²⁰<https://vos.openlinksw.com/owiki/wiki/VOSIntro>

2.6. VIRTUOSO



Figure 2.5: A snapshot of the Virtuoso SPARQL endpoint.

resource directly into the Virtuoso Quad Store. It is also possible to assign a graph IRI where to upload the resource. A snapshot of this feature is shown in Figure 2.6.

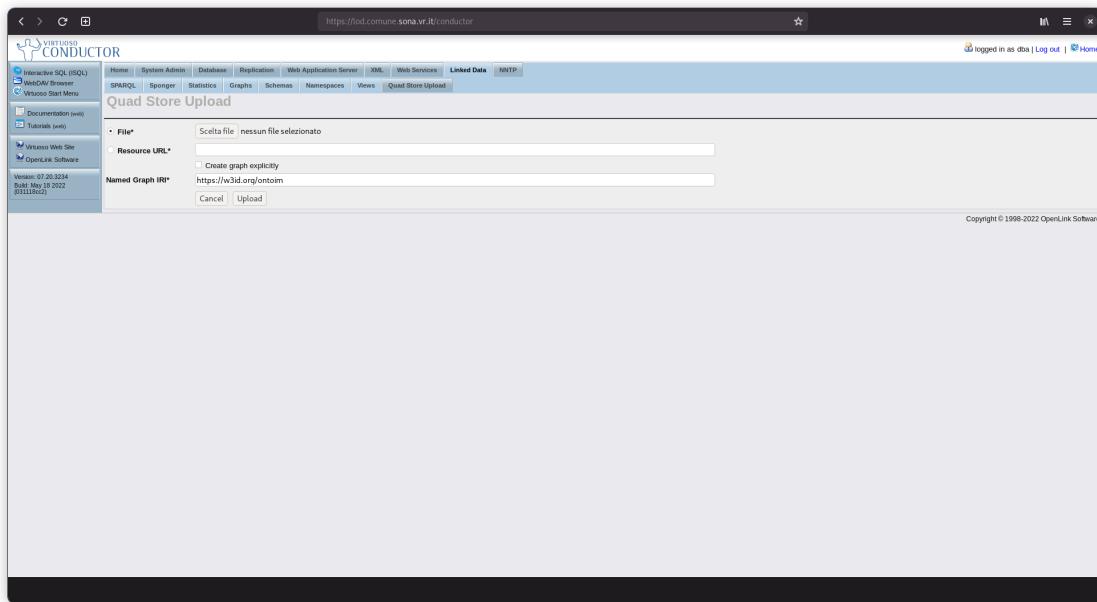


Figure 2.6: A snapshot of the "Quad Store Upload" tab.

RDF Sink Folder WebDAV supports a special folder called `rdf_sink`. This folder

can be used to upload RDF files from any WebDAV client, which are automatically uploaded to the Virtuoso Quad Store.

HTTP PUT RDF files can be uploaded to a `rdf_sink` folder through the HTTP PUT method. Using cURL, an example is:

```
curl -T foaf.rdf
      http://localhost:8890/DAV/home/dba/rdf_sink/foaf.rdf
      -u dba:dba
```

HTTP POST Virtuoso supports HTTP POST method to execute SPARQL/Update language using `Content-Type: application/sparql-query` in the HTTP request headers. Using cURL, an example is:

```
curl -i -d "INSERT {
      <http://w3id.org/people/lucamartinelli>
      <http://www.w3.org/1999/02/22-rdf-syntax-ns#type>
      <http://xmlns.com/foaf/0.1/User>
}" -u "dba:dba"
      -H "Content-Type: application/sparql-query"
      http://localhost:8890/DAV/home/xx/yy
```


SPARQL endpoint If the user has the permission to insert graphs directly from the SPARQL endpoint, using the SPARQL/Update language, as in the example above.

These were just some features provided by Virtuoso Universal Server in order to use it as a SPARQL endpoint and RDF data store system. The full documentation on how to use Virtuoso is available at <https://vos.openlinksw.com/owiki/wiki/VOS>.

2.7. CKAN

2.7 CKAN

The Comprehensive Knowledge Archive Network, or CKAN²¹, is an open source²² data management system. In particular, CKAN is the world's leading tool for making Open Data websites, helping to manage and publish collections of data. It is mainly used by national and local governments, research institutions and other organizations who collect data. Two examples are the U.S. Government's Open Data portal, shown in Figure 2.7a or the Roma Capitale's Open Data portal, shown in Figure 2.7b.



(a) U.S. Government's Open Data portal.^a

^a<https://data.gov/>

(b) Roma Capitale's Open Data portal.^a

^a<https://dati.comune.roma.it/>

Figure 2.7: Examples of CKAN Open Data governments portals.

In CKAN data is published in units called *datasets*. Each dataset is owned by an *organization* and contains information about the data like the title, publisher, data or the license; and one or more resources which are the data itself. For example, a dataset can contain different files, like the data for different years, or the same data in different formats. Any user can view, download, and search for public datasets, but there is also the possibility to restrict the access of some datasets only for registered and authorized users.

Despite the core version of CKAN has only few basic features, one of the strengths of this tool is the possibility to add different plugins which extend its functionalities and customize the user interface. The most popular plugins, developed and maintained by CKAN itself, are: (1) different tools to visualize data directly on the web page, such as tables, plots or maps; (2) DataStore extension

²¹<https://ckan.org/>

²²<https://github.com/ckan/ckan>

that provides an *ad hoc* database for storage of structured data from resources and integrates them into CKAN API to return data in JSON format; (3) DCAT extension that includes RDF serialization of datasets and harvesters to import RDF resources into CKAN. An example of this feature can be seen in the Italian Open Data portal²³, that include the datasets from all the local governments.

2.8 ONTOPIA

The only ontology OntoIM imports is OntoPiA. OntoPiA²⁴ is a network of ontologies and controlled vocabularies developed in 2017 by the Agency for Digital Italy (AgID)²⁵ and the Italian Digital Transformation Team²⁶ with the collaboration of research entities (CNR) and other Italian public administrations (ISTAT, Agenzia delle Entrate, Ministero della Cultura, etc). OntoPiA aims to facilitate the process of data exchange between public administrations, standardize government data, and create the knowledge graph of Italian Public Administration [Dig17b; Dig17a]. Actually the network is composed by 28 ontologies and 39 controlled vocabularies. The OntoPiA ontological stack, shown in Figure 2.8, consists of the following levels:

Foundation Level It's composed by the top-level ontology L_0 , which allows all the ontologies to be linked, enabling the network of ontologies. This ontology defines a few general concepts, such as *Entity*, *Location*, *Activity*, etc, which are used by the ontologies of the upper levels;

Core Level It comprehends the core ontologies, which describes concepts used by different datasets. In particular, the core level describes people, organizations, and locations;

Supporting Level The third level is composed by supporting ontologies, which describe concepts used in the other ontologies. These concepts are: time, roles, measurement units, access conditions, tickets, social media and languages;

²³<https://www.dati.gov.it/>

²⁴<https://github.com/italia/daf-ontologie-vocabolari-controllati/>

²⁵<https://www.agid.gov.it/>

²⁶<https://teamdigitale.governo.it/>

2.8. ONTOPIA

Domain Level The final level comprehends all the ontologies that describe specific domains such as accommodation facilities, events, public contracts, etc.

In addition, there are two metadata ontologies: (1) *DCAT-AP_IT*, an extension of DCAT,²⁷ and DCAT-AP²⁸ ontologies, that aims at facilitating the interoperability between Italian data catalogs; (2) *ADMS-AP_IT*, based on ADMS²⁹, it is used to add metadata to all ontologies in the OntoPiA network. Table 2.3 describes all the ontologies that are part of OntoPiA, with their URIs and prefixes.

Finally, in order to facilitate the interoperability of the data, and let ontology-based application to work properly [EMS08], CPV-AP_IT, CLV-AP_IT, L0-AP_IT, POI-AP_IT, ACCO-AP_IT, Lang-AP_IT, and COV-AP_IT ontologies are aligned with some common ontologies, such as FOAF,³⁰ Org.³¹ or GeoSPARQL³²




Figure 2.8: The OntoPiA ontological stack.

²⁷<https://www.w3.org/TR/vocab-dcat-2/>

²⁸<http://data.europa.eu/r5r/>

²⁹<https://www.w3.org/TR/vocab-adms/>

³⁰<http://xmlns.com/foaf/0.1>

³¹<http://www.w3.org/ns/org#>

³²<http://www.opengis.net/ont/geosparql>

Prefix	URI	Name
Foundation Level		
10	https://w3id.org/italia/onto/l0	Level-0
Core Level		
clvapit	https://w3id.org/italia/onto/CLV	Address (Location)
covapit	https://w3id.org/italia/onto/COV	Organization (Public or Private)
cpvapit	https://w3id.org/italia/onto/CPV	Person
Supporting Level		
acapit	https://w3id.org/italia/onto/AccessCondition	Access Conditions
langapit	https://w3id.org/italia/onto/Language	Language
muapit	https://w3id.org/italia/onto/MU	Value and Measurement Unit
poiapit	https://w3id.org/italia/onto/POI	Points of Interest
potapit	https://w3id.org/italia/onto/POT	Price/Offer/Ticket
roapit	https://w3id.org/italia/onto/R0	Role
smapit	https://w3id.org/italia/onto/SM	Social Media/Contact and Internet
tiapit	https://w3id.org/italia/onto/TI	Time
Domain Level		
accoapit	https://w3id.org/italia/onto/ACCO	Accommodation Facilities
aopapit	https://w3id.org/italia/onto/AtlasOfPaths	Atlas of Paths
chapit	https://w3id.org/italia/onto/CulturalHeritage	Cultural Heritage

2.8. ONTOPIA

Prefix	URI	Name
cis	http://dati.beniculturali.it/cis	Cultural Institute/Site and Cultural Event
cpevapit	https://w3id.org/italia/onto/CPEV	Public Events
cpsvapit	https://w3id.org/italia/onto/CPSV	Public Services
herapit	https://w3id.org/italia/onto/HER	Higher Education and Research
indicator	https://w3id.org/italia/onto/Indicator	Indicator
iotapit	https://w3id.org/italia/onto/IoT	IoT event
parkapit	https://w3id.org/italia/onto/PARK	Parking
pcapit	https://w3id.org/italia/onto/PublicContract	Public Contracts
prjapit	https://w3id.org/italia/onto/Project	Project
rtapit	https://w3id.org/italia/onto/Route	Routes
trapit	https://w3id.org/italia/onto/Transparency	Transparency Obligations
Metadata		
admsapit	https://w3id.org/italia/onto/ADMS	Asset Description Metadata Schema
dcatapit	https://w3id.org/italia/onto/DCAT	Data Catalog Vocabulary

Table 2.3: Ontologies part of the OntoPiA network.

3

Related works

As said in Chapter 1, the purpose of this thesis is to design an ontology for Italian municipalities, facilitate the publication on the Web of Linked Open Data, and develop a web application that makes this data easier for people to comprehend and visualize. It is therefore interesting to understand how major Italian, European, and global cities publish their data on the Web, and what data they publish.

The next sections will show the information collected by Italian, European and global cities about their Open Data portal, and in particular: (1) the number of available datasets; (2) the most common data file types; (3) a score from 1 to 5 based on the five stars classification presented in Section 2.2. Since the score is assigned to the entire data catalog and not to a single resource, only the types of files most present in the portal were considered.

3.1 ITALIAN CITIES

For what concerns Italian cities, has been analyzed the most economically and culturally relevant cities in northern, central, and southern Italy: Bologna, Firenze, Genova, Milano, Napoli, Roma, Torino, and Venezia. The results, collected during April 2022, are shown in Table 3.1. All the cities but Bologna scored three stars, since data are mostly published in non-proprietary format, in particular CSV, JSON, and Shapefile. Firenze and Bologna use API that serves the resources in different formats. Firenze's data can be accessed in JSON format or downloaded

3.1. ITALIAN CITIES

as a ZIP archive containing the CSV file and a metadata file. Bologna reached four stars since it lets export resources in different formats, including RDF/XML, JSON-LD, N-Triples, and Turtle. However, these resources are not accessible through SPARQL, they're not modeled using an ontology, and they're not linked to other data. For these reasons this catalog obtained four stars, and not five.

City	# Datasets	File type	Score	Software
Firenze	1902	Uses API	3	Drupal + CKAN
Bologna	425	Uses API	4	OpenDataSoft
Milano	1618	CSV (1540)	3	CKAN
Torino	1954	CSV (1460)	3	CKAN
Roma	319	CSV (230)	3	CKAN
Venezia	248	CSV (179)	3	Drupal
Genova	138	CSV (111)	3	DKAN
Napoli	62	CSV (35)	3	Custom

Table 3.1: Analysis of Italian cities' Open Data Portals. The data reported in this table was collected during April 2022.

All the Italian cities analyzed, except Napoli, follows the *Linee guida nazionali per la valorizzazione del patrimonio informativo pubblico*.¹ Indeed, they provide their entire catalog as Linked Open Data using the DCAT_AP-IT ontology for resource metadata, like the access and download URL, the name and the file type of the resource, the owner of the dataset, the frequency of updating the data, the theme, and more. This approach aims to maintain the ease of publishing data (e.g. using the CKAN portal), but at the same time allows resources to be more accessible, provide additional information about the nature of the data, and enables the ability to access resources from regional, national,² and European³ portals.

Moving on, an example of nearly five-star data comes from Roma, which provided the list and the information of accommodation facilities⁴ using the OntoPiA ontology described in Section 2.8. However, this data is provided as RDF files, and not using a SPARQL endpoint.

¹<https://docs.italia.it/italia/daf/lg-patrimonio-pubblico/>

²<https://dati.gov.it>

³<https://data.europa.eu>

⁴<https://dati.comune.roma.it/catalog/dataset/suar2021>

Finally, some notable attempts to publish data as Linked Open Data come from Milano and Bologna, which have respectively two portals (Roma⁵, and Bologna⁶) dedicated to Linked Open Data. Milano developed a custom ontology called *OntoMI*⁷ that partially extends OntoPiA, which is described in Section 2.8. In particular, the ontology describes six subject areas that represent a part of the services offered by the City of Milano: libraries, administrative acts, kindergartens, consumer price detection, sports facilities, and Area C entry detection. However, the SPARQL endpoint is no longer available, and the data can no longer be accessed. For what concerns Bologna, it also developed a custom ontology, called *Onto Municipality*⁸, that describes districts, areas, streets, squares and other circulation areas, civic numbering, places and people of interest, schools, and demographic statistics. For the latter, Bologna uses an ontology developed by ISTAT as part of the 2011 census⁹ that is no longer maintained and accessible. Despite the SPARQL endpoint, and the data are still accessible, the project has not been maintained since 2016, making it currently useless.

3.2 EUROPEAN AND GLOBAL CITIES

As for Italian cities, it is interesting to analyze the approach to Open Data (and Linked Open Data) of European and global cities. In particular, has been analyzed the political and economic capitals of major European states, the United States, Canada and Australia. The results, collected during April 2022, are shown in Table 3.2. All the cities except for Amsterdam and London scored three stars, since data are mostly published in non-proprietary format, in particular CSV, JSON, and GeoJSON. On the contrary, Amsterdam and London, despite publishing data under an open license, most resources are available only in the proprietary Excel format. Notice that Berlin, Brussels, Paris, The Hague, New York, Los Angeles, Washington DC, Melbourne, and Sydney use APIs and let export the data in different formats. Brussels, and Paris, which use the same software as Bologna, and New York, Los Angeles, and Melbourne, also allow resources to be exported in RDF format, but without a SPARQL endpoint, without following an ontology,

⁵<https://dati.comune.milano.it/sparql/home.html>

⁶<http://linkeddata.comune.bologna.it>

⁷<https://dati.comune.milano.it/sparql/onthdoc.html>

⁸<http://linkeddata.comune.bologna.it/ontologies/2014/04/onto-municipality/>

⁹<https://www.istat.it/it/archivio/160039>

3.2. EUROPEAN AND GLOBAL CITIES

and without links to other data, so the same considerations about Bologna made in Section 3.1 apply.

City	# Datasets	File type	Score	Software
European cities				
Berlin	2470	Uses API	3	Drupal + CKAN
London	1047	XLSX (644)	2	DataPress
Zurich	683	CSV (463)	3	Custom
Vienna	560	CSV (477)	3	CKAN
Brussels	550	Uses API	4	OpenDataSoft
Barcelona	525	CSV (471)	3	CKAN
Lisbon	359	GeoJSON (206)	3	CKAN
Prague	354	CSV (194)	3	CKAN
Amsterdam	327	XLSX (n.d.)	2	Custom
Paris	321	Uses API	4	OpenDataSoft
The Hague	308	Uses API	3	Dataplateform
Madrid	195	JSON (138)	3	CKAN
Munich	176	CSV (175)	3	CKAN
Global cities				
New York	3541	Uses API	4	Socrata
Los Angeles	1635	Uses API	4	Socrata
Washington DC	1333	Uses API	3	ArcGIS Hub
Toronto	425	CSV (175)	3	WordPress + CKAN
Montreal	320	CSV (227)	3	CKAN
Melbourne	221	Uses API	4	Socrata
Sydney	176	Uses API	3	ArcGIS Hub

Table 3.2: Analysis of European and Global cities' Open Data Portals. The data reported in this table was collected during April 2022.

As for Italian cities, all cities belonging to European Union provides their catalog as Linked Open Data using the DCAT_AP ontology for metadata, in order to make the resources accessible through the European portal.

A similar approach applies to U.S., Canadian and Australian cities, whose catalog of data is collected using the local government API and is made available also in the central government data portal.

To conclude the analysis on Italian, European and global cities' approach to Open Data, we can definitely see that no cities publish Linked Open Data, but they prefer publish resources in using non-proprietary (and in some cases proprietary) format, reaching a score of three or fewer stars. This is probably due to the fact that convert and publishing data as Linked Open Data has a greater cost in terms of time and economic resources [BK11]. The examples of Milano and Bologna, which have stopped investing in Linked Open Data, are proof of this. However, these costs can be covered by states, ministries, government institutions, or regions, which instead publish a portion of their data as Linked Open Data. Some examples are the Europeana project,¹⁰, Ministero della Cultura,¹¹ ISPRA,¹² Regione Veneto,¹³ or Regione Sicilia.¹⁴

Of interest is the approach of Italian and EU cities in publishing the catalog in RDF format using the DCAT metadata profile, which allows them to provide some semantic information to the datasets, such as the owner of the data, the frequency of update, or the topic to which the data refer.

Finally, it is also interesting to analyze the choices of different cities regarding the software chosen to publish Open Data. The most popular tool is CKAN (14 out of 28 cities), especially for Italian and European cities (12 out of 21 cities) use CKAN for their data portal. The reasons for this choice certainly lie in the potential offered by CKAN, which, as explained in Section 2.7, is highly customizable and expandable with plugins, is an open source program and easily installed, and offers the possibility of sharing the catalog in RDF with the DCAT metadata profile, facilitating the interoperability with the central governments.

¹⁰<https://www.europeana.eu>

¹¹<https://dati.cultura.gov.it/>

¹²<http://dati.isprambiente.it/>

¹³<https://www.culturaveneto.it/it/>

¹⁴<https://dati.regione.sicilia.it/i-linked-open-data-nel-catalogo-regionale/>

4

Requirements analysis

As introduced in Chapter 1, this thesis aims to facilitate the publication and dissemination of Open Data, and in particular Linked Open Data, by Italian municipalities by designing and developing an ontology to describe the data, and to develop a web application to make it usable for local government, citizens and businesses to consult the data. In particular, the ontology and the web applications are designed taking in consideration the needs and the data of the Comune di Sona¹, as part of the *Innovation Lab*² project, a project financed by Regione Veneto that aims to spread digital and Open Data culture.

One of the best practices in designing an ontology is to reuse, where possible, existing ontologies [NM+01]. Following this principle, the OntoIM ontology imports the ontologies of the OntoPiA network. As described in Section 2.8, OntoPiA is maintained by AgID and the Italian Digital Transformation Team, and aims to describe different domains of the Italian public administrations, and in particular: people, public and private organizations, addresses and locations, point of interests, accommodation facilities, paths, cultural heritage, cultural events, public services, parking, public contracts, transparency obligations, projects, routes, IoT events, indicators, and higher education and research. Where possible, these ontologies are also aligned with existing ontologies on the web.

The OntoIM ontology was therefore designed and developed as an extension

¹<https://comune.sona.vr.it/>

²<https://innovationlab.regione.veneto.it/>

of the existing OntoPiA ontology, and with the aim that it would become an integral part of the network.

The first part of the design phase involved not only analyzing the data provided by the Comune di Sona, but also analyzing which data the major Italian cities share on their Open Data portals. This choice is due to the fact that we want to create an ontology that can also be reused by other administrations, and takes into account possible future extensions. The work described in Chapter 3, therefore, served not only to analyze how data are made public by various cities, but also what data is available. In addition, some data have been collected from Italian government portals or public agencies, such as Camera di Commercio, ISTAT or Agenzia delle Entrate, to have uniformly structured data across cities.

The data collected comprehends: private organizations, associations, municipal offices, events, cultural heritage, point of interests, accommodation facilities, street directory, traffic and road accidents, municipal heritage, concession acts, waste production, schools and courses organized by private organizations, and demographic statistics (which also includes statistics on tourism, association members, students, and event attendance). In addition to these requests, to make an ontology that is adaptable to other municipalities as well, the census of plants, green areas and street signs, and hospitals were added. Of course, since this is Government Open Data, the privacy of organizations and citizens must also be guaranteed.

Once the data were collected, it was necessary to understand how well OntoPiA ontologies could describe the areas involved. After that, we proceeded to design the ontology by adding the missing classes and properties, and going on to modify the existing ones where necessary.

5

Description of the OntoIM Ontology

The next two sections will describe more in details the OntoIM ontology, designed to describe the semantic areas presented in Chapter 4. In particular, Section 5.1 presents the choices made in developing the ontology, while Section 5.2 will describe the main semantic areas and will present the principal classes in each of them.

5.1 OVERALL DESIGN PRINCIPLES

The design of the ontology started analyzing the data collected from Comune di Sona, and public agencies. Table 5.1 shows what data were collected and where they were collected from.

As said in Chapter 4, the best practice of using existing ontologies where possible was followed. The next step then was to figure out which areas were already described by OntoPiA ontologies and which, instead, needed to be created or imported. The new classes created, moreover, following the design principles of OntoPiA, are subclasses of others existing in OntoPiA ontologies and, in particular, the top-level ontology L0. Indeed, as said in Section 2.8, this ontology allows all the ontologies to be linked, enabling the network of ontologies. The next sections of this thesis will focus on the classes and properties that are strictly part of the OntoIM ontology, while excluding those that are part of the OntoPiA ontology.

The first version of OntoIM ontology is composed of 526 classes, 569 object properties, 405 data properties. The URI of the ontology, the controlled vocabu-

5.1. OVERALL DESIGN PRINCIPLES

Data	Source
Demographic statistics (citizens by location and year, citizenship of foreigners, statistics on names and surnames)	Comune di Sona
Associations	Comune di Sona
Civil status events (births, deaths, emigrations, immigrations, marriages, civil unions, divorces)	Comune di Sona
Concession acts	Comune di Sona
Cultural events	Comune di Sona
List of majors	Comune di Sona
Municipal heritage	Comune di Sona
Museums and cultural heritage	Comune di Sona
Point of Interests	Comune di Sona
Popular University (courses and subscribers)	Comune di Sona
Traffic observations	Local police
Accommodation facilities	Regione Veneto and Comune di Sona
Tourism (arrivals and presences by nationality/region)	Regione Veneto
Private organizations	Camera di Commercio
Addresses and civic numbers	Agenzia delle Entrate
Municipal offices	IPA (AgID)
Waste production	ISPRA
Road accidents	ISTAT
Schools	Ministero dell'Istruzione

Table 5.1: The data collected as reference for designing the OntoIM ontology, and their source.

laries and the resources are secure and permanent by using the W3 Permanent Identifier Community Group, which let create permanent Uniform Resource Locators (URLs) that redirects to defined locations on the Web. Moreover, thanks to this service it was possible to implement a content negotiation mechanisms, to return serialized resources and ontologies in different formats (such as RDF/XML or Turtle), or its visualization/documentation, depending on the request.

The persistent URI, for the OntoIM ontology is <https://w3id.org/ontoim>,

while its prefix is `ontoim`.

Finally, all the files, and the documentation of the ontologies and the controlled vocabularies are open source and available on a GitHub repository¹, which also allows for a permanent location to place the serialization and documentation of the ontology and the other resources.

5.1.1 SEMANTIC AREAS

The OntoIM ontology that extends OntoPiA is divided into nine semantic areas, which describe seven specific domains about the municipality and the local territory. The semantic areas are as follows, and will be explored individually in Section 5.2:

Demographic Observations and Events This semantic area is used to describe not only the number of citizens by year, by geographic area, and by different properties, but also the number of employees that work in an organization, the members part of an association, the number of tourists, the number of subscribers to an event, and so on. It also includes the number of civic status events, like births or deaths, and singular events, like a subscription to an event, to an accommodation facility or a single civil status event, like a marriage;

Facilities and Cadastral Data The entities of this area describe general facilities and their cadastral data;

Organizations and Associations This semantic area extends the CLV_AP-IT ontology of OntoPiA. It describes the private and public organizations, adding information such as the enterprises' life cycle events, the typology of the organizations, and the heritage. This semantic area comprehends also associations, which are treated as private organizations;

Transparency This semantic area includes concession acts and payments from organizations (generally public administrations) to other beneficiaries;

Roads and Traffic The entities of this area describe traffic observations, road signals, and road accidents;

¹<https://github.com/luca-martinelli-09/ontoim>

5.2. AREA-BY-AREA

Schools This semantic area describes public and private schools, comprehensive institutes, and courses organized by public or private organizations;

Green Zones and Plants This area describes green zones, with their information, and the plants with their status;

Hospitals Entities in this area are used to describe hospitals, and hospital departments;

Waste Production It is the area that describes the observations on waste production by year, by waste category and by geographic area.

5.1.2 CONTROLLED VOCABULARIES

For some classes it was necessary to introduce categories or types to classify them. Some examples are the type of civil status event, the category of an association, the type of traffic signal etc. Instead of introducing numerous named individuals into the ontology, some controlled vocabularies and, in particular, taxonomies were chosen. This makes it easier to modify and keep up-to-date the possible categories and types of the various classes, and it was also possible to define hierarchies and subcategories (e.g., a wedding can be religious wedding or civil wedding).

In three cases, named individuals within the ontology were used instead: for traffic direction, for tourist type, and for plant status. Indeed, in these cases, it was not necessary to define a hierarchy, and the types will remain unchanged over time.

To distinguish controlled vocabularies from the entities of the ontology, the URI prefix <https://w3id.org/ontoim/controlled-vocabulary/> has been used.

5.2 AREA-BY-AREA

In this section, each of the semantic areas introduced in Section 5.1.1 will be explored in depth and, in particular, the main elements that comprise it will be described. Bold font will be used to indicate classes, monospaced font for URIs, and italic font for properties. The prefixes in Table 2.3 will be used for URIs.

5.2.1 DEMOGRAPHIC OBSERVATIONS AND EVENTS

This semantic area is represented in Figure 5.1, and comprehends the classes and the properties that describe demographic observations and events for the domains of demographic statistics, civil status, tourism, events, schools, and organizations.

There are two main classes in this area. The **Demographic Observation** class (:DemographicObservation), and the **Demographic Event** (:DemographicEvent) class. This makes it possible to represent data either in aggregate form (e.g., the number of marriages in a year) or, where possible, as individual events (e.g., a wedding occurred on a particular day).

The class **Demographic Observation**, which is a subclass of the **Activity** class (10:Activity) has the data property *observationValue* (xsd:nonNegativeInteger), which represents the quantity of the observation (e.g. the number of citizens). Through the properties *hasTemporalEntity* and *hasSpatialCoverage*, **DemographicObservation** is connected to **TemporalEntity** (tiapit:TemporalEntity) and **Feature** (clvapit:Feature) respectively. In this way it is possible to define the time and space to which the observation refers. To get more accurate statistics, it may be useful to distinguish observations by gender, citizenship or other characteristics. Instead of adding different properties for each of these characteristics, you can use the *hasDemographicReference* property, which connects **DemographicObservation** to a **Person** (cpvapit:Person) class or **Family** (cpvapit:Family) class for observations about families. These classes then reference the observed characteristics.

Different type of observations are modeled using subclasses of the class **DemographicObservation**. The main subclasses are:

Tourists (:Tourists) To describe observations about the number of tourists. We can distinguish two type of observations about tourism: arrivals and presences. The *hasTouristType* property connect **Tourists** to the enumerated class **TouristType** (TouristiType), which can be :Arrival or Presence;

CivilStatus (:CivilStatus) To describe the number of civil status events (e.g. births, deaths, marriages, etc). The typology of the event is defined through the class **CivilStatusCategory** (:CivilStatusCategory), connected by the property *hasCivilStatusCategory*. The available categories are defined in a

5.2. AREA-BY-AREA

controlled vocabulary following the entries in the ISTAT D.7.A model² of civil status statistics.

Subscribers (:Subscribers) This subclass is used to describe observations about subscribers to events, school, and courses.

The remaining subclasses, which doesn't have additional properties, are: **Employees** (:Employees), **Members** (:Members), **Citizens** (:Citizens), and **Bookings** (:Bookings). They describe respectively: (1) the number of employees that works in a organization; (2) the number of members of an association; (3) demographic observation on the population; (4) the number of bookings in a accommodation facility.

The **DemographicEvent** class is similar to **DemographicObservation**, but, as said before, is used to represent singular events. It shares with the latter the *hasDemographicReference* and *hasSpatialCoverage* properties, but it has the property *date*, which defines the date on which the event occurred. Different subclasses are used to model specific types of events. In particular, such subclasses are: (1) **CivilStatusEvent** (:CivilStatusEvent), which is the equivalent of **CivilStats** for the observations; (2) **Subscriber** (:Subscriber), to model a single registration to an event, school or course; (3) **Booking** (:Booking), to describe a single reservation to an accommodation facility.

5.2.2 FACILITIES AND CADASTRAL DATA

Figure 5.2 represent the semantic area of facilities and cadastral data, which are used to describe schools, hospitals, green zones etc.

The main class is **Facility** (:Facility), which is a subclass of the POI-AP_IT's **PointOfInterest** class (poiapit:PointOfInterest), and from which it inherits all the properties. The properties *hasOnlineContactPoint*, *hasPhysicalContactPoint*, and *hasAccessConditions* connect the **Facility** class respectively to: (1) **OnlineContactPoint** (smapit:OnlineContactPoint), which defines online contact points such as emails, social network usernames, websites, and telephones; (2) **PhysicalContactPoint** (smapit:PhysicalContactPoint), which describes the address or the Point of Interest where the facility is located; (3) **AccessCondition**

²<https://purl.archive.org/istat-d7a-sona-2021>

(`acapit:AccessCondition`), which defines the facility opening hours. The **OSDFeature** class (`accoapit:OSDFeature`) is connected through the property *hasOfferedService*, and it is used to describe services and features that the facility offers (e.g. air conditioning, food service, parking, etc).

A facility can be owned by an organization, but can be granted for use to another organization or to a person through a deed of grant. This situation is described by the property *ownedBy*, which connects the facility to the **Organization** class that represents its owner; the property *concessedWithAct* and the class **ConcessionAct** (`:ConcessionAct`), which is described in Section 5.2.4.

Finally, a facility is identified in the land registry by one or more cadastral data. The class **CadastralData** (`:CadastralData`) stores this information, and it is connected to the **Facility** class through the *hasCadastralData* property. A facility has also a cadastral category, defined by the class **CadastralCategory** (`:CadastralCategory`). The elements of this class must be defined according to a controlled vocabulary that stores the available cadastral categories.³

5.2.3 ORGANIZATIONS AND ASSOCIATIONS

This semantic area is represented in Figure 5.3, and covers the domains of public and private organizations, and associations. The area was designed following the structure of the data provided by the Camera di Commercio.⁴

Three data properties concerning events in the life cycle of an organization have been added to the COV-AP_IT's **Organization** class (`covapit:Organization`). These properties, of type `xsd:date`, are *endActivityDate*, *bankruptcyDate*, and *liquidationDate*.

The property *hasLocalUnitAddress* connects the **Organization** class to an **Address** class (`clvapit:Address`), and it is used to store the addresses of one or more local units of the organization. It should also be specified that an organization may have its primary address in another city and a local unit in the municipality concerned. In that case, the Camera di Commercio provides only the address of the latter, specifying that it is a local unit.

An organization may also have real estate assets. Since there can be different

³<https://purl.archive.org/age-categorie-catastali>

⁴<https://www.mn.camcom.gov.it/files/RegistroImprese/Legenda-elenchi.pdf>

5.2. AREA-BY-AREA

types of heritage, such as unavailable heritage or state property, the structures owned by the organization are grouped, through the property *hasFacility*, in the **Heritage** class (:Heritage), to which its type defined by the **HeritageType** class (:HeritageType) is linked, through the property *hasHeritageType*, and whose elements are defined in a controlled vocabulary.

For what concern demographic observations, as specified in Section 5.2.1, the **Employees** class describes the number of employees that works in the organization.

This semantic area also extends the COV-AP_IT's **PrivateOrganization** class (covapit:PrivateOrganization), which is a subclass of an **Organization**. The property *hasOrganizationSection* connect the company to the **OrganizationSection** class (:OrganizationSection). The elements of this class are defined in a controlled vocabulary created from the sections provided by the Camera di Commercio. Another controlled vocabulary defines the elements that must be used for the **CompanyDemographicCategory** class (:CompanyDemographicCategory), which defines whether the enterprise is a youth, female or foreign enterprise. Finally, artisan organizations are described by the class **ArtisanOrganization** (:ArtisanOrganization), which has the two data properties *artisanRegisterCode* (rdfs:Literal), which define the identifier of the organization in the register of artisans, and *artisanRegistrationDate* (xsd:date), which define the date when the organization was registered.

Associations are treated as private organizations. Indeed, the **Association** class (:Association) is a subclass of the **PrivateOrganization** class. The data property *associationRegisterCode* (rdfs:Literal) defines the identifier of the association in the register of associations. The other two data properties, *associationRegistrationDate*, and *associationRemovalFromRegisterDate*, define the life cycle events for an association. The class **AssociationCategory** (:AssociationCategory), connected through the property *hasAssociationCategory*, define the category of the association (such as ONLUS, cultural association, etc). The elements of this class must be defined according to a controlled vocabulary.

Finally, as specified in Section 5.2.1, the **Members** class is used to describe demographic observations on the members of the association, while the **Subscriber** class is used to describe singular membership events.

5.2.4 TRANSPARENCY

This semantic area contains classes that extend the Transparency-AP_IT ontology for transparency obligation. These classes are graphically represented in Figure 5.4.

The main class is **ConcessionAct** (:ConcessionAct), which defines a deed of concession from an organization to another organization or person, and it is used either to grant the management of a facility, or to send a payment. The data properties of the **ConcessionAct** class are: (1) *actTitle* (xsd:string), which defines the title of the act; (2) *actNumber* (xsd:string), which defines the code number of the act; (3) *paymentAmount* (xsd:float), which defines the amount of the payment; (4) *actDate* (xsd:date), which defines the date of the act.

The *hasBeneficiary* property connects the **ConcessionAct** to the **Agent** class (10:Agent), which can be both an organization or a person. The organization that sign the concession act is connected by the *hasOrganization* property.

The signed document of the concession act can be stored as a **TransparencyResource** element (transapit:TransparencyResource).

The referent of the organization that sign the concession act is defined by a **TimeIndexedRole** class (roapit:TimeIndexedRole) through the property *hasActReferent*.

Finally, the typology of the concession act is defined by the **ConcessionAct-Type** class (:ConcessionActType), whose elements are defined in a controlled vocabulary.

5.2.5 ROADS AND TRAFFIC

This semantic area, which is represented in 5.5, extends the OntoPiA ontology regarding the traffics flow observations, and contains classes that describe the road signals and the road accidents.

For what concerns the traffic flow observations, the IOT-AP_IT's **TrafficFlow** class (ioapit:TrafficFlow) is connected through the property *hasTrafficFlowDirection* to the **TrafficFlowDirection** class (:TrafficFlowDirection). This class represents the direction of the flow respect to the sensor, and is thus an enumerated class that can have individual :In or individual :Out as its elements. The **Road-Segment** (:RoadSegment) class, which define the feature of interest for the observation, is connected through the property *hasStreetToponym* to the name of the road

5.2. AREA-BY-AREA

observed, represented by the **StreetToponym** class (`clvapit:StreetToponym`). Finally, the observations can be enriched with the category of vehicle observed. This can be done through the property `hasVehicleCategory` that connects **TrafficFlow** to the **VehicleCategory** class (`:VehicleCategory`). The elements of this class are in a controlled vocabulary, whose elements are defined following the road code classification.⁵

It is also useful for a city to keep track of the status and location of street signs to aid their maintenance. The class **RoadSignal** (`RoadSignal`) describes a street sign, and is a subclass of the **PointOfInterest** class (`poiapit:PointOfInterest`). In addition to those inherited from the **PointOfInterest** class (such as the location), the **RoadSignal** class has three more data properties: (1) `installationDate` (`xsd:date`), which defines the date when the signal was installed; (2) `removalDate` (`xsd:date`), which defines the date when the signal was removed; (3) `signalValue`, which defines the value on the signal (e.g. the speed for a speed limit signal). The typology of the road signal is defined through the property `hasSignalType`, which connects it to the class **RoadSignalType** (`:RoadSignalType`). The elements of this class are defined in a controlled vocabulary generated from the list of road signs in Italy.⁶

The last subarea is dedicated to road accidents, and was developed following the ISTAT model used to collect road accidents, both to design the ontology and to define the controlled vocabularies.⁷

The main class is **RoadAccident** (`:RoadAccident`), which is a subclass of L0's **EventOrSituation** (`l0:EventOrSituation`). The date and time when the incident occurred are defined by the data property `date` (`xsd:dateTime`). The **Geometry** class (`clvapit:Geometry`) describe the geographical location of the accident.

The property `detectedBy` connects **RoadAccident** with the class **ElevationUnit** (`:ElevationUnit`), defining who recorded the incident (e.g., police, carabinieri, etc). The elements of the latter class are defined in a controlled vocabulary. The class **AccidentType** (`:AccidentType`) defines the nature of the accident (e.g. head-on collision, pedestrian investment, etc), and its elements are also defined

⁵<https://www.polizialocaleterredifrontiera.com/codice-della-strada/veicoli/>

⁶https://it.wikipedia.org/wiki/Segnaletica_stradale_in_Italia

⁷<https://purl.archive.org/istat-incidenti-stradali-2018>

in a controlled vocabulary. Another class whose elements are defined in a controlled vocabulary is **WeatherCondition** (:WeatherCondition), which defines the weather conditions at the time of the accident.

The **Road** (:Road) class is used to store information about the road and its status at the time of the accident. The address where the location happened is defined by the **Address** class (clvapit:Address) through the property *hasAddress*. Then, there are several classes that define the context and the status of the road during the accident. The elements of these classes must all be defined according to the relative controlled vocabularies. Such classes are:

RoadCategory (:RoadCategory), which defines whether the road is a municipal road, a state road, a highway, etc;

RoadType (:RoadType), which defines whether the road is single- or multi-lane or one-way;

PavementType (:PavementType), which defines whether the road is paved or uneven;

RoadContext (:RoadContext), which defines whether the accident occurred at an intersection, traffic circle, on a straight road, etc;

RoadbedStatus (:RoadbedStatus), which defines the status of the roadbed (e.g. dry, wet, frozen, etc);

RoadSignalPresence (:RoadSignalPresence), which defines whether the road signals are present or not, or what type of signals are present.

Finally, there are some classes that describe the entities involved (obstacles, persons, vehicles) in the accident. These entities are described by the class **InvolvedEntity** (:InvolvedEntity). The property *hasAccidentCircumstance* connects this class to the **AccidentCircumstance** class (:AccidentCircumstance), which defines what the entity involved was doing during the accident (e.g., the pedestrian was crossing the street, or the car was proceeding at high speed). The elements of this class are defined in a controlled vocabulary. To define whether the entity is a vehicle, an obstacle, or a person, the **InvolvedEntity** class has three subclasses: **InvolvedVehicle** (:InvolvedVehicle), **InvolvedPerson** (:InvolvedPerson), and **InvolvedObstacle** (:InvolvedObstacle). The **InvolvedPerson** class inherits also the properties of the CPV-AP_IT's **Person** class (cpvapit:Person), so it is possible

5.2. AREA-BY-AREA

to define the sex and age of the person involved. The status of the person involved (e.g. wounded, uninjured, dead, etc) in the accident is described using the class **InvolvedPersonStatus** (`:InvolvedPersonStatus`), whose elements are defined in a controlled vocabulary. For what concerns the vehicle, the class **InvolvedVehicle** has the properties *hasFrontPassenger*, *hasBackPassenger*, and *hasConducent*, which are sub-properties of *hasConducent*, and defines the person involved that was inside the vehicle. The **InvolvedVehicle** is also a subclass of the **Vehicle** class (`:Vehicle`). This class, which represents a vehicle, has the data properties *licensePlate* (`xsd:string`), *registrationYear* (`xsd:gYear`), *brand*, *model*, and *color*, which defines the characteristics of the vehicle. The properties *hasEngineDisplacement*, *hasLength*, *hasWidth*, *hasHeight*, and *hasHeight* are used to describe the sizes of the vehicle, using the class **Value** (`muapit:Value`). The category of the vehicle is described by the **VehicleCategory** class. The registration country is defined through the property *hasRegistrationCountry*, which connects the **Vehicle** class to the **Country** class (`clvapit:Country`).

Note that it was still chosen to include properties that in the Open Data environment would not be used for privacy reasons (e.g., the license plate or name of the people involved). This is because in this way the ontology can also be used internally to describe traffic accidents using Linked Data.

5.2.6 SCHOOLS

The classes of this area describes the school, comprehensive institutes, and courses organized by schools. It should also be specified that in this ontology schools also include training institutions, music schools, sports associations, popular universities, etc. Figure 5.6 shows the graphical representation of these classes.

For what concern schools, the class **School** (`:School`) is the core part of the area, and it inherits the properties from the **Facility** class. The data property *schoolCode* represents the MIUR⁸ identifier for the school. The class **SchoolType** (`:SchoolType`) defines the typology of the school (e.g. primary school, high school, etc) and its elements must be defined according to a controlled vocabulary. The subclasses **PublicSchool** (`:PublicSchool`), and **PrivateSchool** (`:PrivateSchool`) are used to distinguish private and public schools. In Italy,

⁸<https://www.miur.gov.it/>

schools are part of comprehensive schools. This situation is described by the **ComprehensiveInstitute** class (`:ComprehensiveInstitute`), and the property *includesSchool*, which links comprehensive institutes to the schools they include.

As said in Section 5.2.1, the classes **Subscribers** and **Subscriber** are used to describe demographic observations on schoolchildren and enrollment events.

For what concern courses, the core class of this subarea is **Course** (`:Course`). The main information of the course are defined by the data properties *courseCode*, which is the identifier of the course, *name*, *description*, *durationHours*, and by the properties: (1) *atTime*, which connects it to a **TimeInterval** class (`:TimeInterval`) to define the period in which the course is provided; (2) *hasPriceSpecification*, which connects **Course** to the class **PriceSpecification** (`:PriceSpecification`) to define the cost of the course, if any; (3) *hasAcademicDiscipline*, which connects the **Course** class to the **AcademicDiscipline** class (`:AcademicDiscipline`) to define the school discipline of the course.

The location (physical or virtual) of the course is specified by the **Classroom** class (`:Classroom`) through the property *situatedInClassroom*. The subclass **PhysicalClassroom** (`:PhysicalClassroom`) defines a physical classroom with an address defined by the class **Address**. On the other hand, if a course is provided using a video communication service, the classroom can be described as a **VirtualClassroom**, specifying also the service used through the class **VideoCommunicationService** (`:VideoCommunicationService`), which describes the name of the service (e.g. Zoom, Google Meet), and its URL.

5.2.7 GREEN ZONES AND PLANTS

Figure 5.7 represent this semantic area, which describes green zones (e.g. parks, flowerbeds, gardens), and plants.

A green zone is treated as a facility. Indeed, the class **GreenZone** (`:GreenZone`) is a subclass of the class **Facility**, from which it inherits the properties. The size of the green zone is defined by the property *hasSurface*, which connects the class to a **Value** class. A green area may have plants within it, so the *hasPlant* property connects a green area with trees within it.

The class **Plant** (`:Plant`) contains the main immutable information about the plant: (1) *plantCode* is the identifier; (2) *commonName* the common name of the

5.2. AREA-BY-AREA

plant; (3) *species* its species; (4) *birthYear* (`xsd:gYear`) the birth year; (5) *plantingDate* (`xsd:dateTime`) the planting date and time. It also inherits the properties from the class **PointOfInterest**, such as the geographical location of the plant. Finally, it is useful to keep track of the state of the plant over time. The class **PlantStatusInTime** (`:PlantStatusInTime`) do this, storing information about the dimensions of the plant, with the properties *hasDiameter*, and *hasHeight*, and the health status, using the enumerated class **PlantHealthStatus** (`:PlantHealthStatus`).

5.2.8 HOSPITALS

Figure 5.8 represents the hospitals.

The main class is **Hospital** (`:Hospital`), which is a subclass of the **Facility** class, and from which it inherits the properties. The data property *totalNumberOfBeds* defines the beds available in the hospital. The two subclasses **PublicHospital** (`:PublicHospital`), and **PrivateHospital** (`:PrivateHospital`) are used to specify the typology of the hospital.

Finally, as for school and comprehensive institutes, the class **HospitalDepartment** (`:HospitalDepartment`) is treated as a hospital and describes the various departments into which a hospital is divided.

5.2.9 WASTE PRODUCTION

This semantic area is represented in Figure 5.9, and defines observations on waste production.

This semantic area is structured as demographic observation. The main class **WasteProduction** (`:WasteProduction`) is connected to a **TemporalEntity** class, and a **Feature** class, which define the space and time to which the observation refers. The value of the observation is defined by the class **Value**, for which you can also define the unit of measurement (e.g. kg or t). To indicate an observation made on only one type of waste, the **WasteCategory** (`:WasteCategory`) class can be used, whose elements are defined in a controlled vocabulary populated with the categories defined by ISPRA.⁹

⁹<https://www.isprambiente.gov.it/it>




Figure 5.1: Demographic Observations and Events semantic area.

5.2. AREA-BY-AREA

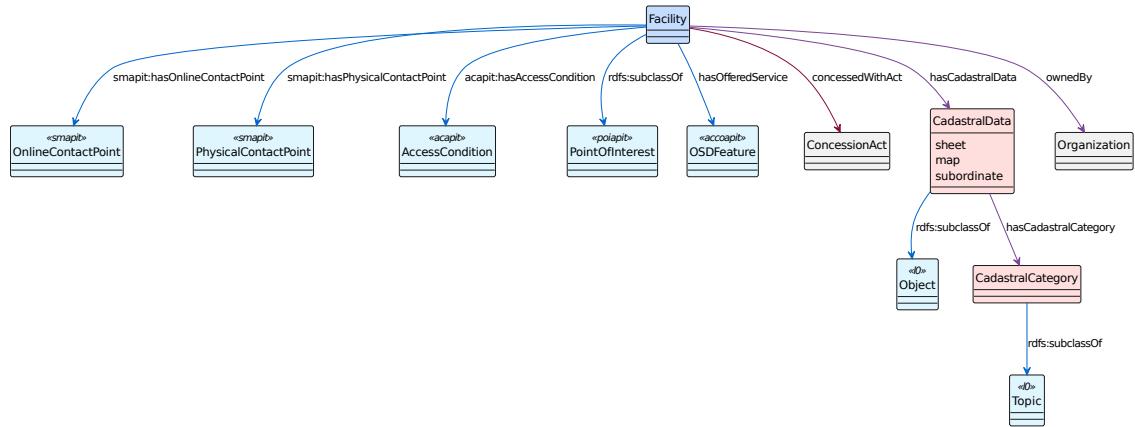


Figure 5.2: Facilities and Cadastral Data semantic area.

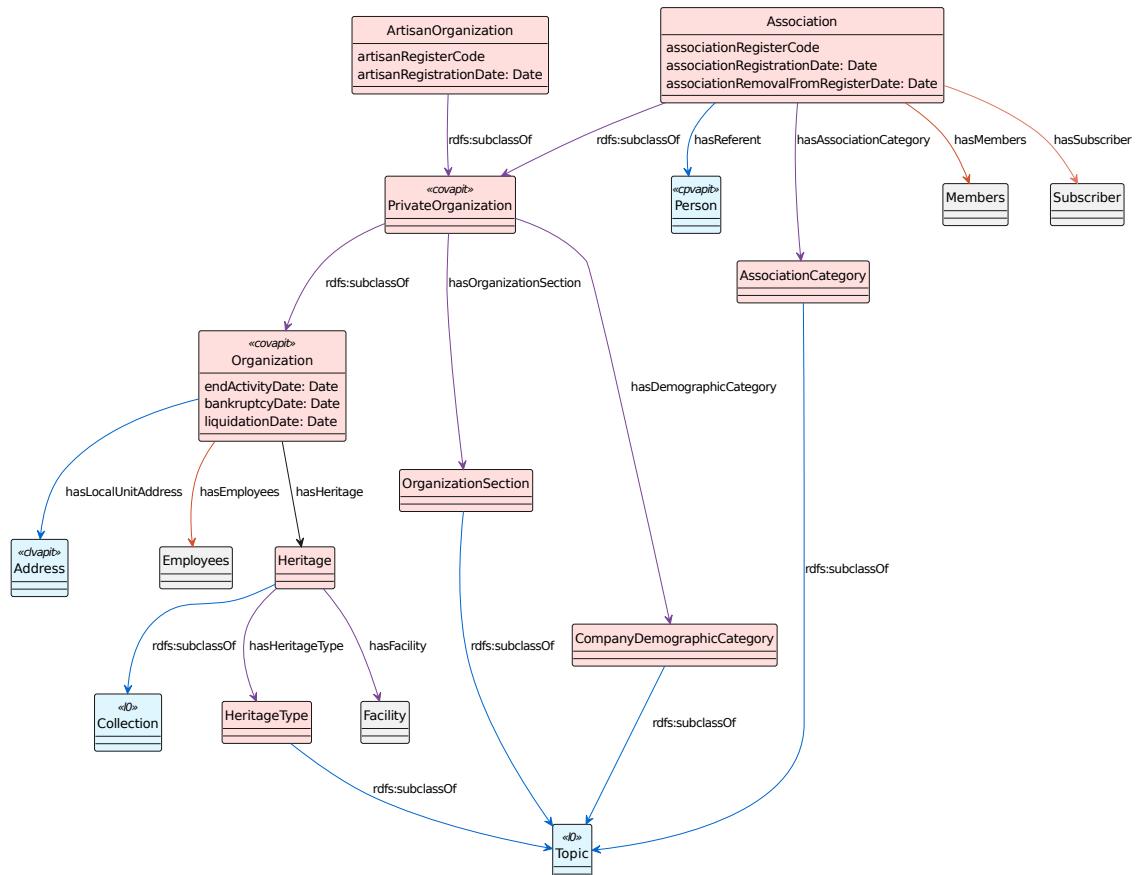


Figure 5.3: Organizations and Associations semantic area.




Figure 5.4: Transparency semantic area.

5.2. AREA-BY-AREA

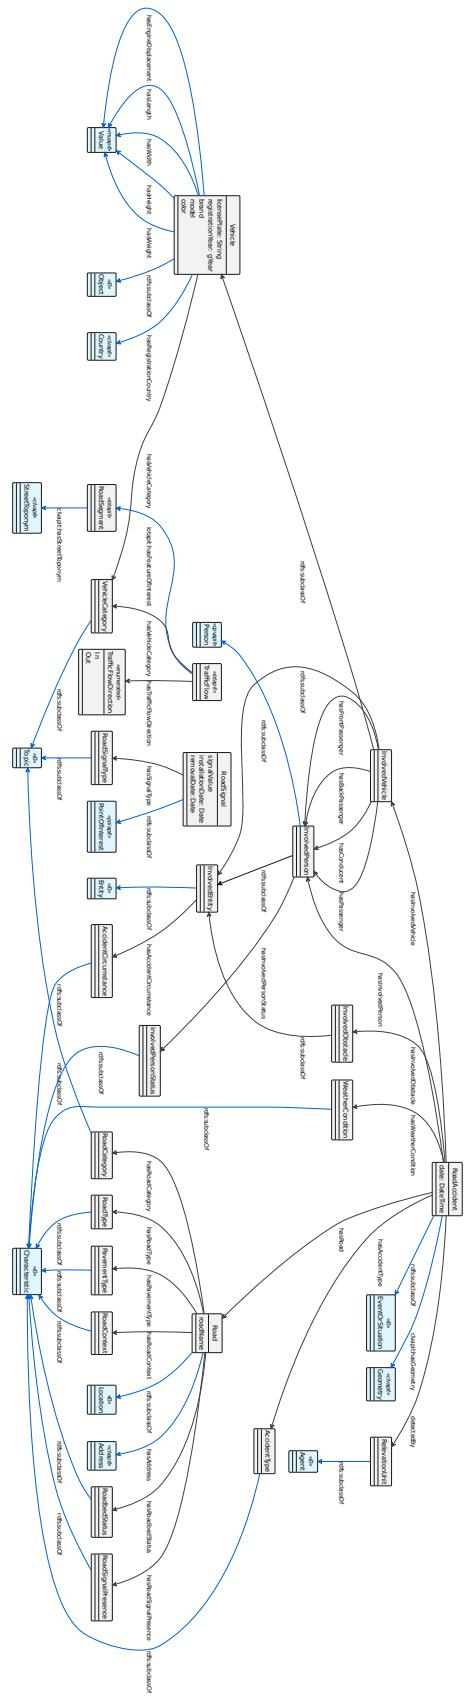


Figure 5.5: Roads and Traffic semantic area.

CHAPTER 5. DESCRIPTION OF THE ONTOIM ONTOLOGY

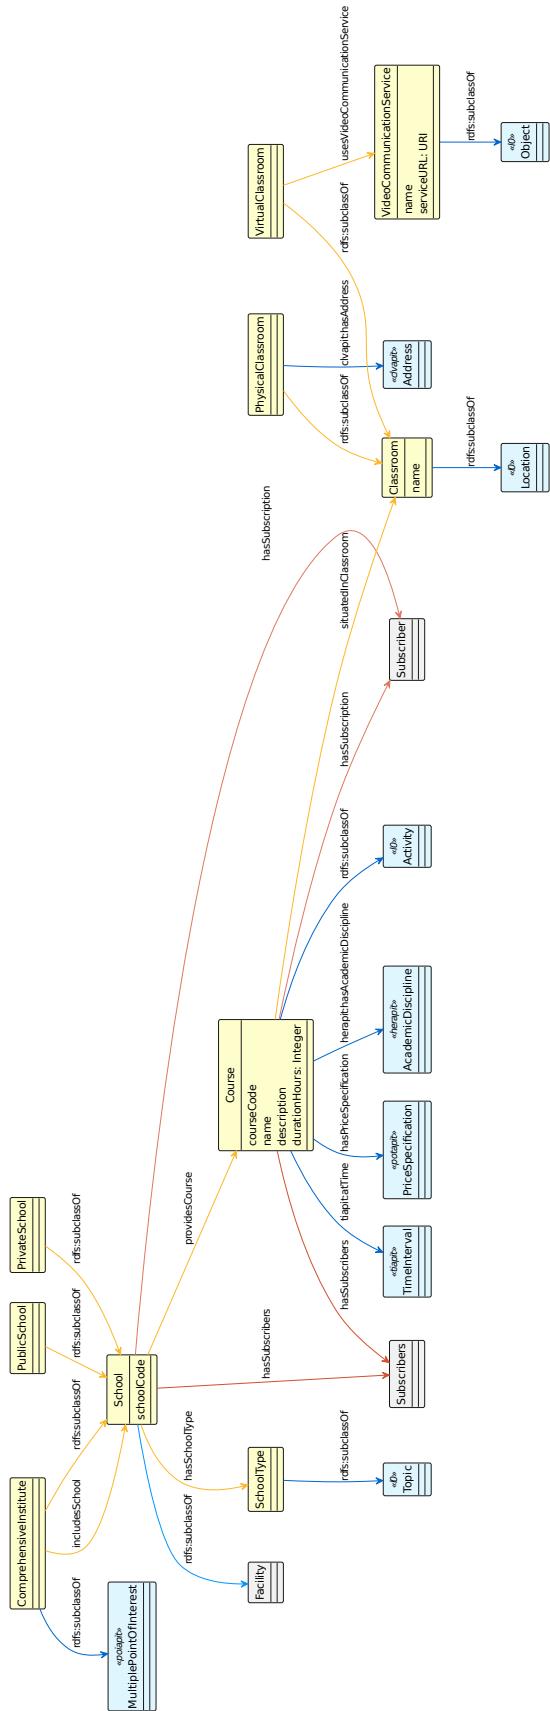


Figure 5.6: Schools semantic area.

5.2. AREA-BY-AREA

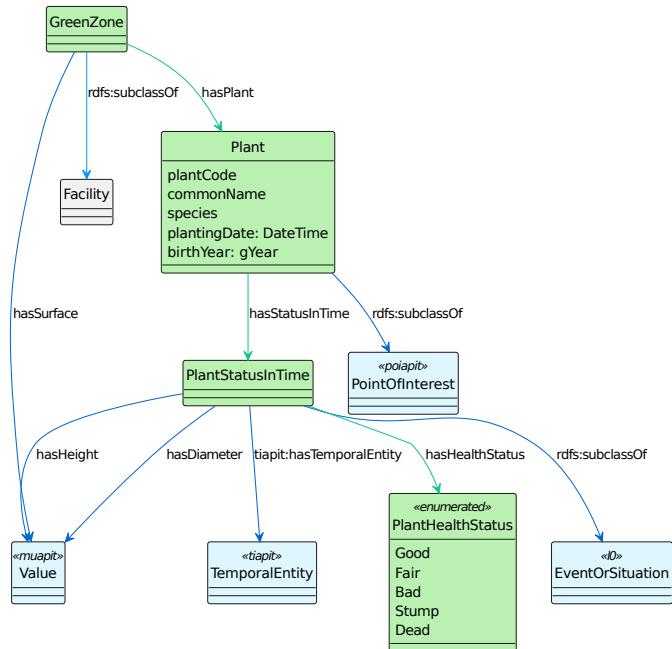


Figure 5.7: Green Zones and Plants semantic area.




Figure 5.8: Hospitals semantic area.

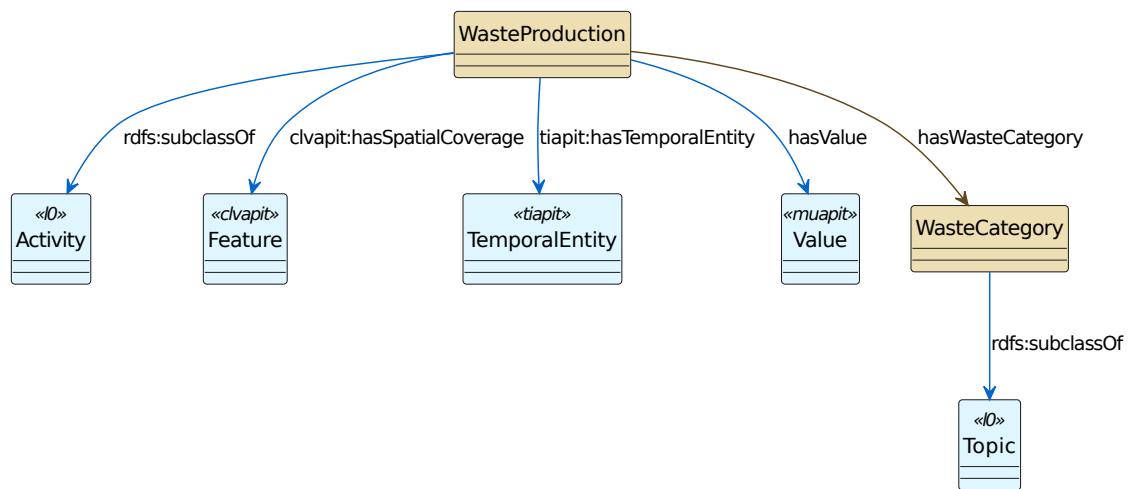


Figure 5.9: Waste Production semantic area.

6

RDF Graph Builder

A fundamental part for providing Linked Open Data on the Web is to build the RDF Graph from the available data. The main obstacle in this project is the fact that different municipalities and public organizations may model the data in different ways, thus making the process of mapping the data into an RDF Graph in accordance with the OntoIM and OntoPiA ontologies difficult. The main idea, as shown in Figure 6.1, is to collect and process data from different sources: the city's CKAN portal, offline files, and Open Data published by state and regional agencies. Thus, two libraries were made in Python to simplify graph creation through object-oriented programming by converting ontology classes into Python classes, and related properties into attributes. This makes data entry from the various sources easier and code reading more understandable. Both the Python libraries and the RDF Graph builder also make use of well-known libraries such as `pandas` and `rdflib` for data manipulation and reading, and graph creation and serialization. Finally, the sources of the various resources and other configurations are all defined in a single `conf.ini` file to allow for better control and easier editing. An example of this file is shown in Code 6.1.

```
1 [ANNCSU]
2 streets = anncsu/streets.csv
3 civics = anncsu/civics.csv
4 census_sections = anncsu/census_sections.kml
5 post_code = 37060
6
7 [ORGANIZATIONS]
8 organizations = organizations/organizations.csv
```

6.1. ONTOPIA-PY AND ONTOIM-PY LIBRARIES

```
9  
10 [WASTE]  
11 waste_production = waste/waste_production.csv
```

Code 6.1: Example of a config.ini file that defines sources for some semantic areas.

Resources will be grouped into datasets defined as dcat:Dataset, and, for the specific case of the Comune di Sona, will be identified by URLs formatted as follows:

https://w3id.org/sona/data/{dataset_id}/{resource_id}

The dataset_id is the identifier of the dataset (e.g. *schools*, *demographic-observations*, etc), while the resource_id is the identifier of the resource, which is generated following different formats depending on the resource. Some supporting resources (e.g., geographic locations or time intervals) may also be located in additional paths that define the resource type (e.g. *geo* or *ti*).




Figure 6.1: RDF Graph Builder architecture.

Section 6.1 will describe the Python libraries developed to facilitate the creation of the RDF Graph, while Section 6.2 will present some examples of data mapping for different data and semantic areas.

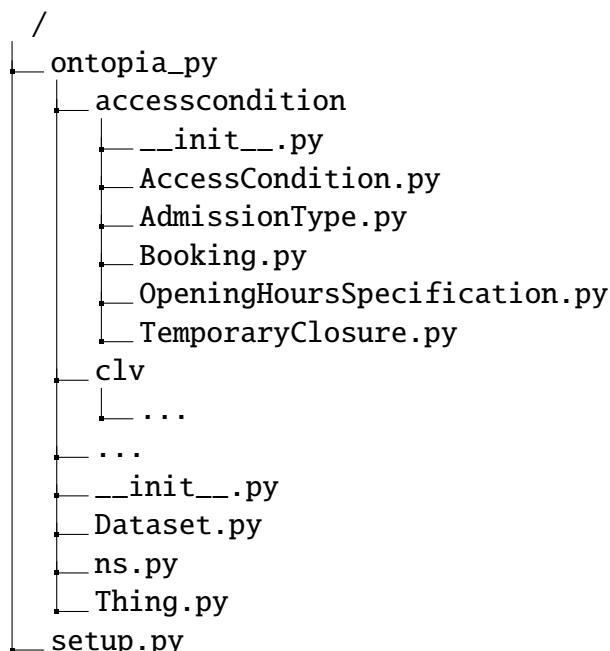
6.1 ONTOPIA-PY AND ONTOIM-PY LIBRARIES

To simplify and facilitate the process of creating the RDF Graph and make the program more understandable, two libraries were created that allow the cre-

ation of resources and the assignment of properties using the object-oriented programming paradigm. These libraries are called `ontopia-py` and `ontoim-py`, for OntoPiA and OntoIM ontologies, respectively. Following the OntoIM ontology, the `ontoim-py` library imports the `ontopia-py` library. The two libraries are also available under an open source license on GitHub¹², and can be installed through the Python PyPI³ package manager, with the command:

```
pip install ontopia_py ontoim_py
```

The `ontopia-py` package is organized as follows, while the `ontoim-py` package follows a similar structure:



This directory `ontopia_py` holds the directories of the OntoPiA's ontologies. Each of these subdirectory holds the files that represent the classes of the ontology. The `ns.py` file contains the namespaces of both ontologies and controlled vocabularies, as shown in Code 6.2.

```

1 from rdflib import Namespace
2
3 L0 = Namespace("https://w3id.org/italia/onto/10/")

```

¹<https://github.com/luca-martinelli-09/ontoim-py>

²<https://github.com/luca-martinelli-09/ontopia-py>

³<https://pypi.org/>

6.1. ONTOPIA-PY AND ONTOIM-PY LIBRARIES

```

4 TRANSP = Namespace("https://w3id.org/italia/onto/Transparency/")
5 TI = Namespace("https://w3id.org/italia/onto/TI/")
6 SM = Namespace("https://w3id.org/italia/onto/SM/")
7 ROUTE = Namespace("https://w3id.org/italia/onto/Route/")
8 RO = Namespace("https://w3id.org/italia/onto/RO/")
9 PUBC = Namespace("https://w3id.org/italia/onto/PublicContract/")
10 PROJ = Namespace("https://w3id.org/italia/onto/Project/")
11 POT = Namespace("https://w3id.org/italia/onto/POT/")
12 POI = Namespace("https://w3id.org/italia/onto/POI/")
13 PARK = Namespace("https://w3id.org/italia/onto/PARK/")
14 MU = Namespace("https://w3id.org/italia/onto/MU/")
15 LANG = Namespace("https://w3id.org/italia/onto/Language/")
16 IOT = Namespace("https://w3id.org/italia/onto/IoT/")
17 INDIC = Namespace("https://w3id.org/italia/onto/Indicator/")
18 HER = Namespace("https://w3id.org/italia/onto/HER/")
19 CULTHER = Namespace("https://w3id.org/italia/onto/CulturalHeritage/")
20 CPV = Namespace("https://w3id.org/italia/onto/CPV/")
21 CPSV = Namespace("https://w3id.org/italia/onto/CPSV/")
22 CPEV = Namespace("https://w3id.org/italia/onto/CPEV/")
23 COV = Namespace("https://w3id.org/italia/onto/COV/")
24 CLV = Namespace("https://w3id.org/italia/onto/CLV/")
25 PATHS = Namespace("https://w3id.org/italia/onto/AtlasOfPaths/")
26 ACOND = Namespace("https://w3id.org/italia/onto/AccessCondition/")
27 ACCO = Namespace("https://w3id.org/italia/onto/ACCO/")
28 ADMS = Namespace("https://w3id.org/italia/onto/ADMS/")
29 CIS = Namespace("http://dati.beniculturali.it/cis/")

```

Code 6.2: Part of the ns.py file that contains the namespaces of OntoPiA's ontologies.

The file `__init__.py` contains the library's main information, like the version number, the description, and the author. Two functions are also defined in this file:

createGraph Initializes the RDF Graph and binds the main namespaces (XSD, SKOS, DCAT, etc) and those declared in the `ns.py` file;

saveGraph Serializes the graph in the different formats.

These two functions are the one showed in Code 6.3.

```

1 def createGraph():
2     g = Graph()
3

```

```

4   g.bind("xsd", XSD)
5   g.bind("foaf", FOAF)
6   g.bind("owl", OWL)
7   g.bind("dc", DC)
8   g.bind("xml", XMLNS)
9   g.bind("dct", DCTERMS)
10  g.bind("rdf", RDF)
11  g.bind("rdfs", RDFS)
12  g.bind("dcat", DCAT)
13  g.bind("prov", PROV)
14  g.bind("skos", SKOS)

15
16  g.bind("l0", L0)
17  g.bind("trapit", TRANSP)
18  g.bind("tiapit", TI)
19  g.bind("smapit", SM)
20  g.bind("rtapit", ROUTE)

21
22  ...
23
24  return g
25
26 def saveGraph(g: Graph, fileName: str):
27     formats = [
28         {"ext": "ttl", "fmt": "turtle"}, 
29         {"ext": "rdf", "fmt": "xml"}]
30     ]
31
32     for format in formats:
33         ext = format["ext"]
34         fmt = format["fmt"]
35
36         with open("{}.{}".format(fileName, ext), "w", encoding="utf-8") as fp:
37             fp.write(g.serialize(format=fmt))

```

Code 6.3: Part of the `__init__.py` file that contains the two functions for creating and saving the graph.

The most important file, that initialize a resource, is `Thing.py`, and all the other classes extends the class `Thing`. The class `Thing` initialize the resource, add the labels and, if specified, connect the resource to a `dcat:Dataset`. In this way, the functions that are responsible for entering the main properties, common to all

6.1. ONTOPIA-PY AND ONTOIM-PY LIBRARIES

resources, are handled by the `Thing` class. The subclasses then only have to deal with defining the attributes, which correspond to the properties to be inserted into the graph.

```
1 class Thing:
2     _dataset: URIRef = None
3     _titles: List[Literal] = []
4     uriRef: URIRef = None
5
6     def __init__(self, id: str, baseUri: Namespace, dataset: Dataset =
7         None, titles: List[Literal] = []):
8         self._dataset = dataset
9         self._titles = titles
10        self.uriRef = URIRef(baseUri[id])
11
12    def _addProperties(self, g: Graph):
13        pass
14
15    def addToGraph(self, g: Graph, isTopConcept=False, onlyProperties=
16        False):
17        if not onlyProperties:
18            g.add((self.uriRef, RDF.type, self.__type__))
19            for title in self._titles:
20                g.add((self.uriRef, DC.title, title))
21
22            if self._dataset:
23                g.add((self.uriRef, SKOS.inScheme, self._dataset.uriRef))
24
25            if isTopConcept:
26                g.add((self._dataset.uriRef, SKOS.hasTopConcept, self.uriRef))
27
28        self._addProperties(g)
```

Code 6.4: The ontopia-py's `Thing` class.

Finally, the `rdfs:subClassOf` property is handled through object-oriented programming inheritance. In the Code 6.5 example, the `Sequence` class, defined in the `L0` ontology, inherits the properties of the `Collection` class, which inherits those of the `Entity` class, which is a subclass of `Thing`. This way it is not necessary to consult the ontology documentation to trace back all the properties that a given class admits, and it is more difficult to make mistakes in creating the graph. In addition, the class type, which is mapped to the `rdf:type` property, is defined by each class in the `__type__` attribute, which defines the URI of the ontology class.

```

1 # Sequence.py
2 class Sequence(Collection):
3     __type__ = L0["Sequence"]
4
5     hasFirstMember: Entity = None
6     hasLastMember: Entity = None
7
8     def _addProperties(self, g: Graph):
9         super().__addProperties(g)
10
11     if self.hasFirstMember:
12         g.add((self.uriRef, L0["hasFirstMember"], self.hasFirstMember.
13             uriRef))
14
15     if self.hasLastMember:
16         g.add((self.uriRef, L0["hasLastMember"], self.hasLastMember.
17             uriRef))
18
19 # Collection.py
20 class Collection(Entity):
21     __type__ = L0["Collection"]
22
23     hasMember: List[Entity] = None
24
25     def _addProperties(self, g: Graph):
26         super().__addProperties(g)
27
28     if self.hasMember:
29         for hasMember in self.hasMember:
30             g.add((self.uriRef, L0["hasMember"], hasMember.uriRef))

```

Code 6.5: The ontopia-py's Sequence and Collection classes. Thanks to the object-oriented programming inheritance it is possible to map ontology classes and properties into Python classes and attributes.

Regarding ontoim-py library, the structure is the same as that of ontopia-py. The `ns.py` file imports all the OntoPiA's namespaces, and declares the OntoIM one. The `__init__.py` file imports the ontopia-py's `saveGraph` function, and extends the `createGraph` function adding the new bindings. Classes such as `Organization`, whose declaration is shown in Code 6.6, for example, which extend the OntoPiA ontology class, inherit attributes and functions from the ontopia-py library, adding the new ones declared in OntoIM.

6.1. ONTOPIA-PY AND ONTOIM-PY LIBRARIES

```
1 from ontopia_py.cov.Organization import Organization
2
3 class Organization(Organization):
4     __type__ = ONTOIM["Organization"]
5
6     hasEmployees: List[Employees] = None
7     hasHeritage: List[Heritage] = None
8     hasLocalUnitAddress: List[Address] = None
9     endActivityDate: Literal = None
10    liquidationDate: Literal = None
11    bankruptcyDate: Literal = None
12
13    def _addProperties(self, g: Graph):
14        super()._addProperties(g)
15
16        if self.hasEmployees:
17            for hasEmployees in self.hasEmployees:
18                g.add((self.uriRef, ONTOIM["hasEmployees"], hasEmployees.
19                      uriRef))
20
21        if self.hasHeritage:
22            for hasHeritage in self.hasHeritage:
23                g.add((self.uriRef, ONTOIM["hasHeritage"], hasHeritage.uriRef))
24
25        if self.hasLocalUnitAddress:
26            for hasLocalUnitAddress in self.hasLocalUnitAddress:
27                g.add((self.uriRef, ONTOIM["hasLocalUnitAddress"],
28                      hasLocalUnitAddress.uriRef))
29
30        if self.endActivityDate:
31            g.add((self.uriRef, ONTOIM["endActivityDate"], self.
32                  endActivityDate))
33
34        if self.liquidationDate:
35            g.add((self.uriRef, ONTOIM["liquidationDate"], self.
36                  liquidationDate))
37
38        if self.bankruptcyDate:
39            g.add((self.uriRef, ONTOIM["bankruptcyDate"], self.bankruptcyDate
40))
```

Code 6.6: The ontopim-py's Organization class. This class inherits the one defined in ontopia-py, declaring the new attributes.

The process of creating and saving an RDF Graph using these two libraries is the following one:

1. Initialize the graph with the function `createGraph`;
2. Create a `dcat:Dataset` object using the class `Dataset`, specifying its URI and adding properties like the name or the author of the Dataset;
3. Add the dataset to the graph calling the function `addToGraph`;
4. Initialize the object that must be inserted into the graph, specifying the identifier of the resource, the base URI, the dataset into which it is to be inserted, and the name of the resource;
5. Add the properties to the resource declaring the attributes of the object created in the step before;
6. Add the resource to the graph calling the function `addToGraph`;
7. Save the graph using the function `saveGraph`, specifying the filename where to save it.

The example code of Code 6.7 creates an RDF Graph with the *ANNCSU* dataset, into which a resource of type `StreetToponym` is inserted.

```

1 from rdflib import XSD, Graph, Literal, Namespace
2
3 from ontopia_py import Dataset, createGraph
4 from ontopia_py.clv import StreetToponym
5
6 # Set namespace for data
7 DATA: Namespace = Namespace("https://w3id.org/sona/data/")
8 ANNCSU: Namespace = Namespace("https://w3id.org/sona/data/ANNCSU/")
9
10 # Create the graph and bind the namespace
11 g = createGraph()
12 g.bind("anncsu", ANNCSU)
13
14 # Create the concept scheme
15 ANNCSU_DATASET: Dataset = Dataset(DATA["ANNCSU"])
16 ANNCSU_DATASET.label = [
17     Literal("Numeri civici e strade urbane", lang="it"),
18     Literal("Civic Addressing and Street Naming", lang="en")
19 ]
20

```

6.2. DATA MAPPING FOR DIFFERENT SEMANTIC AREAS

```
21 # Add to graph
22 ANNCSU_DATASET.addToGraph(g)
23
24 # Create the street toponym
25 streetToponym: StreetToponym = StreetToponym(
26     id="street-1",
27     baseUri=ANNCSU,
28     dataset=ANNCSU_DATASET,
29     titles=[Literal("Via Roma", datatype=XSD.string)]
30 )
31 streetToponym.toponymQualifier = "Via"
32 streetToponym.officialStreetName = "Roma"
33
34 # Add to graph
35 streetToponym.addToGraph(g)
```

Code 6.7: An example of the creation of an RDF Graph with the ontoim-py and ontopia-py libraries.

6.2 DATA MAPPING FOR DIFFERENT SEMANTIC AREAS

After presenting and describing the ontopia-py and ontoim-py libraries, the next sections will show how the data entry of the Comune di Sona took place for the following semantic areas: Addresses, Organizations, and Schools. Indeed, these three semantic area are respectively examples of data mapping from the CKAN Open Data Portal, offline file, and government Open Data portals.

All written programs make use of common functions grouped in a package called `utils`. These functions are: (1) `getConfig`, to retrieve the configuration file; (2) `getOpenData`, to retrieve data from the three different sources presented in Chapter 6 and return it in the form of pandas `DataFrame`; (3) `standardizeName`, which transforms strings into a standard format where only initials are capitalized, and apostrophe characters are converted to the relevant accented characters; (4) `genNameForID`, which generates an identifier from the name, removing special characters and using a hyphen as a separator. This function is used only in some cases where the name is known to be unique, such as for localities. In other cases, it is preferable to use identifiers such as vat number or a sequence number.

The full code of the RDF Graph Builder for the Comune di Sona is available on GitHub at <https://github.com/luca-martinelli-09/sona-lod>.

```
1 def getConfig(fileName):
2     config = configparser.ConfigParser()
3     config.read(fileName)
4
5     return config
6
7 def getOpenData(resID, baseURL=None, whereSQL="", rawData=False, dtype=None, strip=True):
8     config = getConfig('../..../conf.ini')
9
10    offline = False if baseURL else config.getboolean("API", "use_offline")
11
12    baseURL = baseURL if baseURL else config.get("API", "base_url")
13
14    dataURI = "{}/api/3/action/datastore_search_sql?sql=SELECT * FROM
15        \"{}\" {}".format(baseURL, resID, whereSQL) if not offline else "
16        ..../off_data/{}".format(resID)
17
18
19    if resID.startswith("http"):
20        offline = False
21        dataURI = resID
22
23
24    if rawData:
25        if offline:
26            return dataURI
27
28        getDataRequest = urlopen(dataURI)
29
30        return getDataRequest
31
32    df = None
33
34    if offline:
35        df = pd.read_csv(dataURI, dtype=dtype)
36    else:
37        if resID.endswith("csv"):
38            df = pd.read_csv(dataURI, dtype=dtype)
39        else:
40            tries = 0
41            res = {"success": False}
42            while not res["success"] and tries < 20:
```

6.2. DATA MAPPING FOR DIFFERENT SEMANTIC AREAS

```
38     res = requests.get(dataURI).json()
39     if res["success"]:
40         df = pd.DataFrame(res["result"]["records"], dtype=dtype)
41         break
42     tries += 1
43
44 if strip and not df is None:
45     df = df.applymap(lambda x: x.strip() if type(x) == str else x)
46
47 return df
48
49 def standardizeName(name):
50     name = name.lower().title()
51
52     if name.endswith("a'"):
53         name = name.removesuffix("a'") + "à"
54
55     return name.strip()
56
57 def genNameForID(name):
58     nameID = ""
59
60     name.replace("'", "")
61     name = unidecode.unidecode(name.lower())
62
63     for char in name:
64         nameID += char if char.isalnum() else "-"
65
66 return nameID
```

Code 6.8: The common functions in the utils package.

6.2.1 ADDRESSES

The addresses (streets and civic numbers) is the only semantic area that relies exclusively on the OntoPiA ontology. In particular, the address ontology is CLV_AP-IT, and was created in collaboration with ISTAT and the Agenzia delle Entrate. The ontology was developed to describe the data collected in the National Archive of Urban Street Numbers, or ANNCSU⁴, a computerized repository

⁴<https://purl.archive.org/anncsu>

containing the street and house numbers of all Italian municipalities. Access to the archive should be public, but at the moment it is possible to download a municipality's data only through the dedicated portal of the Agenzia delle Entrate and reserved for the municipal contact person for toponymy, who is responsible for managing the updating of the data. From this dedicated portal it is possible to download the list of streets and the list of house numbers. The files are in CSV format, and follows a specific structure designed by the two agencies⁵. The information about the census sections are available on the ISTAT website, as kml file that defines the geographic geometries of each section.

The URL of the ANNCSU Dataset for the Comune di Sona is <https://w3id.org/sona/data/ANNCSU/>, so the resources will be organized in this Dataset and that URL will be their namespace. In detail:

Sona Is the municipality, for which `owl:sameAs` properties have been defined that link it to external resources concerning the city, such as dbpedia, ISPRA, and the Ministry of Cultural Heritage. The URL of this resource is <https://w3id.org/sona/data/ANNCSU/023083>, where 023083 is the ISTAT identifier for the city;

Census Sections The file of the census sections has been processed thanks to the pykml Python library, thanks to which it is possible to retrieve the section number and its geographic geometry. The URL of the resources is formatted as <https://w3id.org/sona/data/ANNCSU/cs/{id}>, where `id` is the number of the census section. The URL of the geometry resource that defines the geographical polygon of the census section is instead formatted as <https://w3id.org/sona/data/ANNCSU/gsc/{id}>;

Streets Names The street names are retrieved from the file of the streets provided by the Agenzia delle Entrate. The URL of the resources is formatted as https://w3id.org/sona/data/ANNCSU/street/{id_street}, where the `id_street` is the national identifier for the street, which can be retrieved by the field PROGR_NAZIONALE in the CSV file;

Civic Numbers As for street names, the civic numbers are retrieved from the file of house numbers provided by the Agenzia delle Entrate. The resources are

⁵<https://purl.archive.org/age-anncsu-specifiche>

6.2. DATA MAPPING FOR DIFFERENT SEMANTIC AREAS

located at https://w3id.org/sona/data/ANNCSU/civic/{id_num}, where `id_num` in this case is the field `PROGR_CIVICO`, and it is the national identifier of the house number;

Addresses Addresses are combination of street names and civic numbers. These are the resources part of the ANNCSU dataset, and their URLs are formatted as https://w3id.org/sona/data/ANNCSU/ad-{id_street}-{id_num}. For each street, in addition to house numbers, an address with `id_num` equals to `snc` is created, indicating the absence of a house number. Finally, for some addresses, thanks to OpenStreetMap⁶, it was possible to estimate the location, which is described in geometry resources located at https://w3id.org/sona/data/ANNCSU/gcn/{id_num}.

Code 6.9 shows the part of the `config.ini` file relative to the ANNCSU, and the part of the RDF Graph Builder that inserts the streets toponyms into the graph.

```
1 # config.ini
2 [ANNCSU]
3 streets = 7b3401e5-7235-43d4-b8ed-8e11f30ec404
4 civics = e22af300-c934-4d7e-b84a-3f58e0862e4c
5 census_sections = 1f929641-0aa5-49b7-b012-6ea29ceca759
6 post_code = 37060
7
8 # rdf_builder_anncsu.py
9 config = getConfig("conf.ini")
10
11 anncsuAddresses = getOpenData(config.get("ANNCSU", "streets"))
12 anncsuAddresses.set_index("PROGR_NAZIONALE", inplace=True)
13
14 for streetID, address in anncsuAddresses.iterrows():
15     dugName = standardizeName(address["DUG"])
16     streetName = standardizeName(address["DENOM_COMPLETA"])
17
18     fullName = "{} {}".format(
19         standardizeName(address["DUG"]),
20         standardizeName(address["DENOM_COMPLETA"]))
21
22     streetToponym = StreetToponym(
23         id="street/" + str(streetID),
```

⁶<https://www.openstreetmap.org>

```

24     baseUri=ANNCSU ,
25     dataset=ANNCSU_DATASET ,
26     titles=[Literal(fullName, datatype=XSD.string)])
27
28     streetToponym.toponymQualifier = [Literal(dugName, datatype=XSD.
29         string)]
30     streetToponym.officialStreetName = [Literal(streetName, datatype=XSD.
31         string)]
32
33     streetToponym.addToGraph(g)

```

Code 6.9: The part of the RDF Graph Builder that inserts the streets toponyms into the graph, and the config.ini file relative to the addresses.

6.2.2 ORGANIZATIONS

As specified in Table 5.1, the information on private organizations comes from the Camera di Commercio. Some of this information, for privacy reasons, such as the tax code, cannot be published. However, thanks to the tax code it is possible to categorize businesses into women's, youth and foreign businesses, as specified in Section 5.2.3. It is therefore necessary, in this case, to retrieve the data from an offline file that contains all the information, and publish the censored version on Open Data portals.

The URL of the dataset of organizations situated in the Comune di Sona is <https://w3id.org/sona/data/organization/>, so the resources will be organized in this Dataset and that URL will be their namespace. Resources are identified by their vat number, which is unique to each enterprise. Only in the case of individuals for whom no vat number is indicated, the tax number is used.

In addition to enterprises, instances describing the online contact points of organizations have as their URL https://w3id.org/sona/data/organization/ocp/{vat_num}, which in turn link back to pecs located at URL https://w3id.org/sona/data/organization/pec/{id_pec}, where id_pec is generated with the genNameForID function described in Section 6.2.

As said in Section 5.2.3, the Camera di Commercio provides a list of all the organizations present in the city, both if they are local units and if they are the head office. The typology is specified in the CSV file in the field UL-SEDE. All the main offices of the organizations are then entered first. Next, the enterprises indicated as local units are processed and, if already present in the graph, the

6.2. DATA MAPPING FOR DIFFERENT SEMANTIC AREAS

`hasLocalUnitAddress` property is added along with their address. Otherwise, the organization is inserted without specifying the `hasPrimaryAddress` property, but only `hasLocalUnitAddress`.

Another problem encountered at this stage concerns the addresses of organizations, but also of other resources in other semantic areas, such as associations, schools, and municipal offices. The addresses, in fact, are present in text format and, often, do not correspond to the official addresses in the national archive (e.g., "Via Francesco Petrarca" is listed as "Via Petrarca") and do not have standard formatting. This is due to the fact that the address is given by the company and is not normalized by the Camera di Commercio. The goal is to retrieve from each address the street identifier and that of the house number (if any). To do this, the program makes use of the function shown in Code 6.10. This function retrieves the street directory from the files presented in Section 6.2.1, and compares the address to be searched against this list, returning a ranking of the most similar results. The `extract` function of the `rapidfuzz` library is used for this part. From here, the desired address must be manually selected, and the two searched identifiers will be returned.

```
1 def queryStreetCode(q):
2     config = getConfig('../..../conf.ini')
3
4     streetsDF = getOpenData(config.get("ANNCSU", "streets")).set_index(["PROGR_NAZIONALE"])
5     civicsDF = getOpenData(config.get("ANNCSU", "civics")).set_index(["PROGR_CIVICO"])
6
7     streetsForSearchIDs = [(c["PROGR_NAZIONALE"], progrCivico) for
8         progrCivico, c in civicsDF.iterrows()]
9     streetsForSearch = ["{} {} {}{} {}".format(
10         streetsDF.loc[c["PROGR_NAZIONALE"]]["DUG"],
11         streetsDF.loc[c["PROGR_NAZIONALE"]]["DENOM_COMPLETA"],
12         c["CIVICO"],
13         "" if pd.isna(c["ESPONENTE"]) else c["ESPONENTE"],
14         streetsDF.loc[c["PROGR_NAZIONALE"]]["LOCALITA'"],
15         ).lower() for _, c in civicsDF.iterrows()]
16
17     streetsForSearchIDs.extend([(progrNazionale, None) for progrNazionale,
18         _, in streetsDF.iterrows()])
19     streetsForSearch.extend(["{} {} {}".format(s["DUG"], s["DENOM_COMPLETA"],
20         s["LOCALITA'"]).lower() for _, s in streetsDF.
```

```

18     iterrows())
19
20     searchResults = process.extract(q.lower(), streetsForSearch, scorer=
21         fuzz.WRatio, limit=10)
22
23     print(f"\n\n[RESULTS FOR] {q}")
24     for res, val, i in searchResults:
25         print(f"[{i}] {res} ({val})")
26
27     selectedResult = input("Choose one or type custom search: ")
28
29     if selectedResult.isnumeric():
30         return streetsForSearchIDs[int(selectedResult)]
31     elif selectedResult == "":
32         return None, None
33     else:
34         return queryStreetCode(selectedResult)

```

Code 6.10: The function that retrieve the street identifiers from the string of the address.

The function is then used as follows:

```

1 address = organizationInfo["INDIRIZZO"]
2 progrNazionale, progrCivico = queryStreetCode(address) if not pd.isna(
    address) else (None, None)

```

6.2.3 SCHOOLS

The list of schools present in the municipality is directly taken from the official data portal of the Ministero dell'Istruzione⁷. In this way, as with businesses, the data structure is standard nationwide, and the code can be easily reused without making further changes. From the portal are available both public and private schools, and the relative comprehensive institutes, for which the comprehensive institute code is the same as the school code. Other information regards the address, the email, and the website of the schools.

The URL of the dataset of schools situated in the Comune di Sona is <https://w3id.org/sona/data/school/>, so the resources will be organized in this Dataset and that URL will be their namespace. Resources are identified by the code

⁷<https://dati.istruzione.it/opendata/>

6.2. DATA MAPPING FOR DIFFERENT SEMANTIC AREAS

assigned by the Ministero dell’Istruzione, which is unique to each school. As for organizations, the online contact point are located at https://w3id.org/sona/data/school/ocp/{id_school}, while the email, the certified email, and the website are located at https://w3id.org/sona/data/school/{contact_type}/{id_c}, where `contact_type` is respectively `email`, `pec`, and `web`, and `id_c` is an identifier generated by the `genNameForID` function.

The same address considerations made for organizations in Section 6.2.2 also apply.

Finally, the typology of the school are mapped using a Python dictionary into the relatively identifiers of the controlled vocabulary described in Section 5.2.6

For what concerns demographic observations on the schoolchildren, the Ministero dell’Istruzione provides, for the public schools, information about the number of the children, both categorized by gender, by course, and by age.

In this case, the URLs of the observations resources are defined as follows: <https://w3id.org/sona/data/school/statistics/{id}>, where `id` is the sequence number of the data inserted. The demographic references are located at <https://w3id.org/sona/data/demographic-references/{id}>. In this case the `id` can be `M` for males, `F` for females, or `age-{age}` for the age. For what concerns the temporal entity that defines the period to which the observation refers, the resources of type `TemporalEntity` are located at <https://w3id.org/sona/data/ti/{sy}-{ey}>, where `sy` stands for the beginning year of the academic year, and `ey` for the end. Since in this case the temporal entity represents a period, the start date is fixed to `{sy}-09-01`, and the end date to `{ey}-07-01`.

Code 6.11 shows the part of the code that build the RDF Graph for demographic observations on the schools’ schoolchildren.

```
1 for (schoolCode, academicYear), statsInfo in sumDataDF.iterrows():
2     academicYear = int(str(academicYear)[:4])
3
4     school = School(
5         id=schoolCode,
6         baseUri=SCHOOL_DATA
7     )
8
9     schoolInfo = schoolsDF.loc[schoolCode]
10
11    temporalEntity = TimeInterval(
12        id="ti/{}-{}".format(academicYear, academicYear + 1),
```

```

13     baseUri=SCHOOL_DATA ,
14     dataset=SCHOOL_DATASET ,
15     titles=[Literal("{} / {}".format(academicYear, academicYear + 1),
16     datatype=XSD.string)]
17 )
18 temporalEntity.startTime = Literal(str(academicYear) + "-09-01",
19     datatype=XSD.date)
20 temporalEntity.endTime = Literal(str(academicYear + 1) + "-07-01",
21     datatype=XSD.date)
22 temporalEntity.addToGraph(g)
23
24 for sexCode in ["M", "F"]:
25     subscribers = Subscribers(
26         id="statistics/{}/{}/{}".format(schoolCode, sexCode, academicYear
27     ),
28         baseUri=SCHOOL_DATA ,
29         dataset=SCHOOL_DATASET ,
30         titles=[
31             Literal("{} - {} - A.A. {} / {}".format(
32                 "Alunni (maschi)" if sexCode == "M" else "Alunne (femmine)",
33                 standardizeName(schoolInfo["DENOMINAZIONE_SCUOLA"]),
34                 academicYear,
35                 academicYear + 1
36             ), datatype=XSD.string)
37         ]
38     )
39
40     demReference = AlivePerson(
41         id="demographic-reference/" + sexCode ,
42         baseUri=SCHOOL_DATA ,
43         dataset=SCHOOL_DATASET ,
44         titles=[
45             Literal("Male" if sexCode == "M" else "Female", lang="en"),
46             Literal("Maschio" if sexCode == "M" else "Femmina", lang="it")
47         ]
48     )
49     demReference.hasSex = Sex(id=sexCode , baseUri=PERSON_SEX)
50     demReference.addToGraph(g)
51
52     subscribers.hasDemographicReference = demReference
53     subscribers.hasTemporalEntity = temporalEntity
54
55     subscribers.observationValue = Literal(statsInfo["ALUNNI"] + ("
```

6.2. DATA MAPPING FOR DIFFERENT SEMANTIC AREAS

```
51    "GENDER": "MASCHI" if sexCode == "M" else "FEMMINA")], datatype=XSD.  
52    positiveInteger)  
53  
54    subscribers.addToGraph(g)  
55  
56    school.hasSubscribers = [subscribers]  
57    school.addToGraph(g, onlyProperties=True)
```

Code 6.11: The part of the code that build the RDF Graph for demographic observations on the school.

7

Web Applications

The third and final key part of this thesis project concerns the web applications that allow on the one hand to download and, for the municipality, upload Open Data, and on the other hand to visualize Linked Open Data clearly through summary tables and graphs. The first web application is CKAN, and it is described in Section 7.1. In this case, the work was to configure the open source portal CKAN, described in Section 7.1, with the necessary plugins and extensions to follow the *Linee Guida Nazionali per la Valorizzazione del Patrimonio Informativo Pubblico*.¹ The second web application, described in Section 7.2, was built from scratch and is intended to retrieve data from the SPARQL endpoint and display it graphically in the form of graphs, maps and tables.

7.1 CKAN

As said in Chapter 6, the RDF Graph Builder get the data from three different sources, and one of these sources is the CKAN Open Data Portal of the city (Comune di Sona in this case). The main idea is to use this open source portal in order to facilitate the publication of data on the web and be aligned with other Italian and foreign cities. At the same time, this portal provides the catalog compliant with the DCAT-AP_IT metadata profile, and this allows published data to be made available in regional, national and European portals as well. In

¹<https://docs.italia.it/italia/daf/lg-patrimonio-pubblico/it/stabile>

7.1. CKAN

addition, the presence of data in Linked Open Data format does not preclude the use of the CKAN portal, despite the fact that the latter's data achieve a three-star rating according to Tim Berners-Lee's classification introduced in Section 2.2. In fact, in addition to having a lower management cost [BK11], this portal allows the publication of data that in Linked Open Data format would not be publishable, or for which an ontology describing them has not yet been developed. It also allows original resources to be reused for other tasks, either by the municipality or by companies or citizens.

The entire project is available under an open source license on GitHub at the link <https://github.com/luca-martinelli-09/ckan>. The final result built for the Comune di Sona is shown in Figure 7.1. For ease of installation and deployment, moreover, the portal has been containerized to produce a Docker² image.




Figure 7.1: A snapshot of the Comune di Sona's CKAN Open Data portal.

The Open Data portal extends the CKAN 2.9 Docker image provided by the Open Knowledge Foundation,³ installing the required plugins. Such plugins are:

²<https://www.docker.com/>

³<https://github.com/okfn/docker-ckan>

View, DataStore, and DataPusher These are the default plugins provided by CKAN. The View plugin is used to show a preview of the resources in form of tables or images. DataStore and DataPusher are two plugins that process CSV files and enable the ability to use APIs to access data in JSON format. Thanks to these plugins, it is also possible to filter data using the SQL language;

Spatial This plugin adds geospatial capabilities to CKAN, adding a spatial field to the dataset schema;

GeoView This extension is used together with the spatial plugin, and contains view plugins to display geospatial files in CKAN;

Harvest It provides a harvesting framework that allows data and catalogs to be imported from external APIs, DCAT catalogs, or other CKAN portals;

Multilang This extension provides a way to localize the CKAN's title and description contents for Dataset, Resources, Tags, Organizations, and Groups;

DCAT and DCAT-AP_IT These two plugins expose the data catalog as RDF document serialized using the DCAT and DCAT-AP_IT profiles.

To further simplify the process of setting up the portal, there is a .env configuration file that allows the main portal information to be specified, such as the account and password and administrator, the name, description and logo of the site, the email and vat number of the public administration that manages the portal, and a possible configuration of the email server for sending notifications. Code 7.1 shows the configuration template file.

```

1 CKAN_SITE_URL # URL of the custom CKAN's portal
2
3 CKAN_SYSADMIN_NAME # Name of the sysadmin account
4 CKAN_SYSADMIN_PASSWORD # Password for the sysadmin account
5 CKAN_SYSADMIN_EMAIL # Email for the sysadmin account
6
7 CKAN_ORG_VAT # IPA/VAT/IVA of the organization
8 CKAN_ORG_EMAIL # Email of the organization
9
10 CKAN__DATAPUSHER__CALLBACK_URL_BASE # Same as CKAN_SITE_URL
11
12 # Core

```

7.2. DATA REPORTS

```
13 CKAN__SITE_TITLE # The website title (usually the P.A. name)
14 CKAN__SITE_DESCRIPTION # Website description
15 CKAN__SITE_ABOUT # Website long description
16 CKAN__SITE_LOGO # Url for the logo
17 CKAN__FAVICON # Url for the favicon
18
19 # Email setup
20 CKAN_SMTP_SERVER
21 CKAN_SMTP_STARTTLS
22 CKAN_SMTP_USER
23 CKAN_SMTP_PASSWORD
24 CKAN_SMTP_MAIL_FROM
25
26 CKANEXT__DCAT__BASE_URI # Same as CKAN_SITE_URL
27
28 GEONAMES__USERNAME # Follow the guide for geosolutions-it/ckanext-
    dcatapit extension
```

Code 7.1: Configuration template file for the CKAN Open Data portal.

Finally, a theme was developed for CKAN, which is installed on par with an extension, to apply graphical customization to the portal. The theme is also available under an open source license on GitHub at <https://github.com/luca-martinelli-09/ckanext-cdsttheme/>.

7.2 DATA REPORTS

The second web application, called *Data Reports*, aims to retrieve, through preset SPARQL queries, the Linked Open Data uploaded in the phase described in Chapter 6, process them and display them in an understandable way by means of tables, graphs and maps. In fact, the basic idea is to create an online journal where different data are displayed and commented on for each article and grouped by semantic area.

As for the CKAN portal described in Section 7.1, the application is containerized in a Docker image to simplify its installation and adoption. The application is open source, and it is available on GitHub at <https://github.com/luca-martinelli-09/ontoim-webapp>. A snapshot of the web application's home page configured for the Comune di Sona is shown in Figure 7.2.

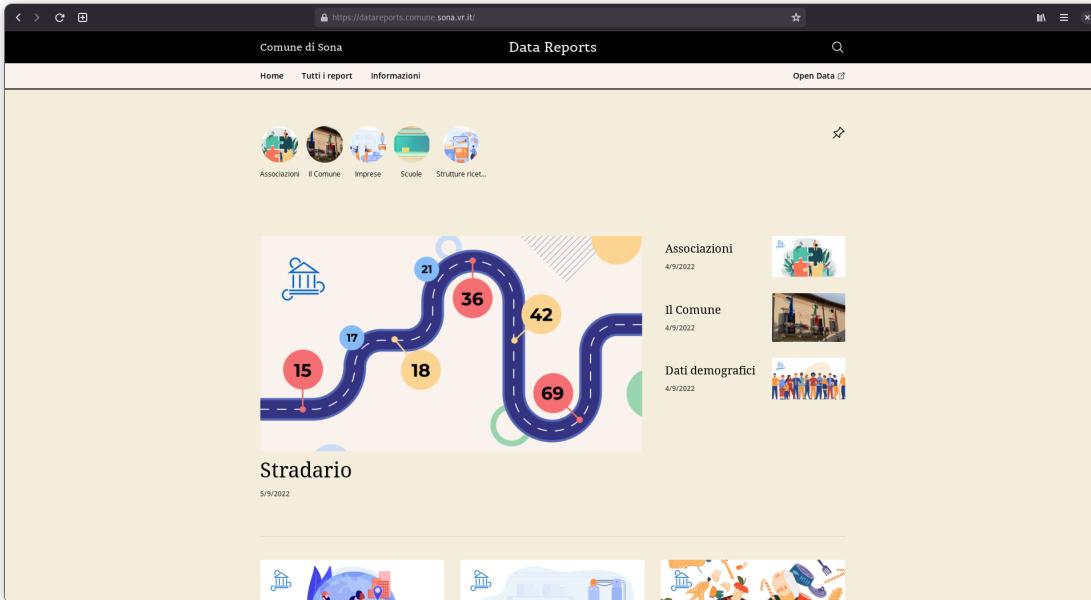


Figure 7.2: A snapshot of the Data Reports home page.

Data Reports was developed using Svelte,⁴ and SvelteKit,⁵ a modern JavaScript framework for building fast and SEO friendly web applications. Unlike other frameworks such as Vue⁶ or React,⁷ in fact, Svelte produces all the files necessary for the site to function at the *compile step*. In addition, Svelte, being a JavaScript framework, can be integrated with numerous ready-made packages available on the npm⁸ package manager. Svelte also supports the creation of components, which are small pieces of code that can be imported and reused on multiple pages. Some of these packages are called *preprocessors*, which Svelte uses to compile code and produce web pages. Some of these, for example, create SEO friendly links, others compact HTML code, and others generate CSS. The full documentation of all the features of Svelte is available at <https://kit.svelte.dev/docs/introduction>.

Data Reports uses several packages and preprocessors, the most relevant of which are:

TailwindCSS and PostCSS They provide a modern CSS framework that aims to

⁴<https://svelte.dev/>

⁵<https://kit.svelte.dev/>

⁶<https://vuejs.org/>

⁷<https://reactjs.org/>

⁸<https://www.npmjs.com/>

7.2. DATA REPORTS

speed up website development, easy to customize, adaptable, and using CSS best practices;

MDsveX It is a preprocessor for Svelte that process Markdown files and convert them to HTML. It also supports Svelte components, scripts, HTML tags, and YAML headers, which let add information to the pages, like the publication date or the title;

Apache ECharts It is an open source library for data visualization. It allows you to easily create interactive charts, maps, and more;

Leaflet It is the leading open source library for creating interactive maps using OpenStreetMap⁹.


Each page/report is therefore a Markdown file that contains a YAML header which specifies the title of the report, the publishing date, a URL for the thumbnail image, and some tags used for classification. Then, the text (e.g. description of the data, some comments on them), is written in Markdown in the body of the page. Code 7.2 shows the Markdown page that displays the list of associations registered in the Comune di Sona, the map where these associations are located, and a pie chart on the number of associations by type. Figure 7.3 shows the pages related to demographic statistics and citizenship of foreign citizens. To increment the readability and reusability of the code, each part of the page involved in retrieving and displaying data was converted into a Svelte component that could be embedded.

```
1 ---
2 title: Associazioni
3 date: 2022-09-04
4 fixed: true
5 thumb: /reports/thumb-associations.png
6 keywords:
7   - associazioni
8 ---
9
10 <script>
11   import TabellaAssociazioni from "../data/associazioni/
12     TabellaAssociazioni.svelte";
```

⁹<https://www.openstreetmap.org/>

```

12 import MappaAssociazioni from "../data/associazioni/MappaAssociazioni
13   .svelte";
14 import TipologiaAssociazioni from "../data/associazioni/
15   TipologiaAssociazioni.svelte";
16 </script>
17
18 In questa sezione è possibile consultare l'albo delle associazioni del
19   Comune di Sona, con le relative informazioni.
20 <TabellaAssociazioni />
21
22 Di seguito la mappa delle associazioni
23 <MappaAssociazioni />
24
25 ## Tipologia di associazioni
26 <TipologiaAssociazioni />
```

Code 7.2: "Associazioni" report page in Markdown.

(a) "Dati demografici" report. (b) "Flussi migratori e cittadinanza" report.

Figure 7.3: Snapshots for different Data Reports's pages.

Each of these components is responsible for running a SPARQL query to the endpoint, and processing the result, in JSON format, to constitute the configuration that generates the charts, tables and maps. The function that execute the SPARQL query, called `querySPARQL`, is common to all components, so it was defined in a `utils` script that is imported. Code 7.3 shows the `querySPARQL` function.

```

1 export const querySPARQL = async (options) => {
2   const urlParams = new URLSearchParams({
3     infer: true,
4     sameAs: true,
5     format: "application/sparql-results+json",
6     query: options.query,
```

7.2. DATA REPORTS

```

7   })
8
9   const response = await fetch(import.meta.env.VITE_SPARQL_ENDPOINT + ""
10    ?" + urlParams.toString());
11
12  if (response.ok) {
13    let res = await response.json()
14    options.success(options.raw ? res : res.results.bindings)
15  } else {
16    if (options.error) {
17      options.error(response)
18    }
19 }

```

Code 7.3: The querySPARQL function.

Instead, Code 7.4 shows the Svelte component that produces the sunburst chart visible in Figure 7.3a about the number of citizens by locality and gender.

```

1 <script>
2   const query = '
3     prefix tiapit: <https://w3id.org/italia/onto/TI/>
4     prefix clvapit: <https://w3id.org/italia/onto/CLV/>
5     prefix ontoim: <https://w3id.org/ontoim/>
6     prefix cpvapit: <https://w3id.org/italia/onto/CPV/>
7     prefix sex: <https://w3id.org/italia/controlled-vocabulary/
8       classifications-for-people/sex/>
9     prefix dc: <http://purl.org/dc/elements/1.1/>
10
11     select ?Localita (sum(?Maschi) as ?Maschi) (sum(?Femmme) as ?Femmme) (sum(?Totale) as ?Totale) where {
12       {
13         select ?te ?spatialCoverage ?Maschi where {
14           ?popM tiapit:hasTemporalEntity ?te ;
15             clvapit:hasSpatialCoverage ?spatialCoverage ;
16             ontoim:hasDemographicReference ?df ;
17             ontoim:observationValue ?Maschi .
18             ?df cpvapit:hasSex sex:M .
19       }
20     } union {
21       select ?te ?spatialCoverage ?Femmme where {
22         ?popF tiapit:hasTemporalEntity ?te ;
23             clvapit:hasSpatialCoverage ?spatialCoverage ;
24             ontoim:hasDemographicReference ?df ;
25       }
26     }
27   }
28 }
```

```

24         ontoim:observationValue ?Femmmine .
25         ?df cvapit:hasSex sex:F .
26     }
27   } .
28   ?te tiapit:year "2021"^^xsd:gYear .
29   ?spatialCoverage dc:title ?Localita .
30
31   bind ((coalesce(?Maschi, 0) + coalesce(?Femmmine, 0)) as ?Totale)
32 } group by ?Localita order by asc(?Localita)';
33
34 let data;
35
36 querySPARQL({
37   query: query,
38   success: (res) => {
39     data = sparqlToArray(res);
40   }
41 });
42 </script>
43
44 {#if data}
45   <Graph
46     option={{{
47       series: {
48         type: "sunburst",
49         data: data
50         .map((el) => {
51           return {
52             name: el["Localita"],
53             value: parseInt(el.Totale),
54             children: [
55               { name: "Maschi", value: parseInt(el.Maschi) },
56               { name: "Femmmine", value: parseInt(el.Femmmine) },
57             ],
58           };
59         }),
60       },
61     }}}
62   />
63 {/if}

```

Code 7.4: The sunburst chart about the number of citizens by locality and gender.

Finally, as for the CKAN Open Data portal, a configuration file is used to

7.2. DATA REPORTS

contain the main constants such as the SPARQL endpoint URL, and the main information about the municipality. Code 7.5 shows this .env configuration file.

```
1 VITE_SPARQL_ENDPOINT # SPARQL endpoint URL
2 VITE_PA_NAME # Public Administration name
3 VITE_PA_URL # Public Administration website
4 VITE_PA_CF # Public Administration tax code
5 VITE_PA_ADDRESS # Public Administration address
6 VITE_OPEN_DATA_URL = # CKAN Open Data portal URL
7 VITE_LODVIEW_HOMEURL = # Eventual LodView main page URL
```

Code 7.5: The DataReports configuration file.

8

Conclusions and Future Works

The goal of this thesis is to document the design and the development of the first version of the OntoIM ontology, the RDF Graph Builder, and the *Data Reports* and CKAN web applications. The goal is to simplify and encourage the adoption by Italian municipalities of technologies to publish data produced by the public administration in Linked Open Data format, and make it more understandable for citizens, businesses, and the administration itself to consult this data.

The OntoIM ontology was designed as an extension of OntoPiA, the Italian ontology developed by AgID, Italian Digital Transformation Team, in cooperation with other national agencies. In particular, the OntoIM was designed in order to describe the data collected from the Comune di Sona, Agenzia delle Entrate, Camera di Commercio, ISPRA, ISTAT, Ministero dell'Istruzione, and after an analysis on the data published by Italian, and global cities. The result is an ontology that describes different domains of a city, such as demographic observations, schools, hospitals, waste production, private organizations, associations, etc., so that it is possible to describe as many aspects of a city as possible, and to be able to publish that data in Linked Open Data format. The first version of the OntoIM ontology, and its documentation, can be found at <https://w3id.org/ontoim>.

The second part of the project was the development of two Python libraries, `ontopia-py`,¹ and `ontoim-py`.² These libraries allow an RDF Graph to be created using the object-oriented programming paradigm, converting ontology classes

¹<https://github.com/luca-martinelli-09/ontopia-py>

²<https://github.com/luca-martinelli-09/ontoim-py>

to Python classes, and properties to attributes. This makes it possible to insert resources more easily, resulting in more readable code, and following ontology rules. These libraries were used to build the RDF Graph of the Comune di Sona by importing data from three different sources: the CKAN Open Data portal, from offline files, and government APIs. The RDF Builder for the Comune di Sona is available on GitHub at <https://github.com/luca-martinelli-09/sona-lod>.

The third and final part of the project was the development of two web applications. The first one is the configuration of a CKAN Open Data portal for the Italian municipalities that follows the *Linee guida nazionali per la valorizzazione del patrimonio informativo pubblico*.³ The second one is *Data Reports* a web application built from scratch, using the Svelte JavaScript framework. *Data Reports* has been designed as an online journal where each article can contain text, and charts, maps, and tables built using the data retrieved by SPARQL queries. The goal is to have a portal where citizens, businesses and the administration can consult data in an easy and understandable way, with comments on the data and results. The two projects are available respectively at <https://github.com/luca-martinelli-09/ckan>, and <https://github.com/luca-martinelli-09/ontoim-webapps>.

As future works, the OntoIM ontology will be subdivided into different ontologies, one for each semantic area, to follow the principle of modularity of the Italian OntoPiA ontology, and with the goal of becoming part of it.

For what concerns the RDF Builder, the future project is to develop a universal and easy-to-setup mapper to convert data from different sources into Linked Open Data, according to the OntoIM and OntoPiA ontologies.

Finally, the *Data Reports* web application will be improved so that it can be maintained through a private area that allows adding and creating charts directly from the web interface.

³<https://docs.italia.it/italia/daf/lg-patrimonio-pubblico/>

References

- [Gru95] Thomas R Gruber. "Toward principles for the design of ontologies used for knowledge sharing?" In: *International journal of human-computer studies* 43.5-6 (1995), pp. 907–928.
- [BHL01] Tim Berners-Lee, James Hendler, and Ora Lassila. "The semantic web". In: *Scientific american* 284.5 (2001), pp. 34–43.
- [NM+01] Natalya F Noy, Deborah L McGuinness, et al. *Ontology development 101: A guide to creating your first ontology*. 2001.
- [Gen+03] John H Gennari et al. "The evolution of Protégé: an environment for knowledge-based systems development". In: *International Journal of Human-computer studies* 58.1 (2003), pp. 89–123.
- [Bec+04] Sean Bechhofer et al. "OWL web ontology language reference". In: *W3C recommendation* 10.2 (2004), pp. 1–53.
- [Ber06] Tim Berners-Lee. "Linked Data - Design Issues". In: (July 2006). URL: <http://www.w3.org/DesignIssues/LinkedData.html>.
- [Fei+07] Lee Feigenbaum et al. "The semantic web in action". In: *Scientific American* 297.6 (2007), pp. 90–97.
- [EMS08] Jérôme Euzenat, Adrian Mocan, and François Scharffe. "Ontology alignments". In: *Ontology Management*. Springer, 2008, pp. 177–206.
- [Hit+09] Pascal Hitzler et al. "OWL 2 web ontology language primer". In: *W3C recommendation* 27.1 (2009), p. 123.
- [Tay10] Mohammad Mustafa Taye. "Understanding semantic web and ontologies: Theory and applications". In: *arXiv preprint arXiv:1006.4567* (2010).
- [BK11] Florian Bauer and Martin Kaltenböck. "Linked Open Data: The essentials". In: *Edition mono/monochrom, Vienna* 710 (2011).

REFERENCES

- [BHB11] Christian Bizer, Tom Heath, and Tim Berners-Lee. "Linked data: The story so far". In: *Semantic services, interoperability and web applications: emerging concepts*. IGI global, 2011, pp. 205–227.
- [Con+13] World Wide Web Consortium et al. "SPARQL 1.1 Overview". In: (Mar. 2013). URL: <https://www.w3.org/TR/sparql11-overview/>.
- [ISA13] ISA Programme. *How Linked Data is transforming eGovernment*. 2013. URL: <https://joinup.ec.europa.eu/node/61404>.
- [Con+14a] World Wide Web Consortium et al. *RDF 1.1 Concepts and Abstract Syntax*. Feb. 2014. URL: <https://www.w3.org/TR/rdf11-concepts/>.
- [Con+14b] World Wide Web Consortium et al. "RDF 1.1 Primer". In: (June 2014). URL: <https://www.w3.org/TR/rdf11-primer/>.
- [IB14] Mirjana Ivanović and Zoran Budimac. "An overview of ontologies and data resources in medical domains". In: *Expert Systems with Applications* 41.11 (2014), pp. 5158–5166.
- [Dig17a] Agenzia per l'Italia Digitale. "Linee guida nazionali per la valorizzazione del patrimonio informativo pubblico". In: (2017). URL: <https://docs.italia.it/italia/daf/lg-patrimonio-pubblico/>.
- [Dig17b] Agenzia per l'Italia Digitale. "Piano triennale per l'informatica nella Pubblica amministrazione 2017-2019". In: (2017). URL: <https://docs.italia.it/italia/piano-triennale-ict/pianotriennale-ict-doc/it/2017-2019/>.

Acknowledgments

First, I would like to thank my parents, Fabiana and Renato, who have guided me and allowed me to get this far, supporting me both emotionally and financially throughout my growth.

I also want to thank my friends, who have always been by my side during these years, with whom I have shared unforgettable experiences, and who have put up with me even and especially during the most difficult times, and without whom I probably would not be the person I am now.

I would also like to thank Prof. Silvello, who believed in this project from the very beginning, and guided me to enable me to bring it to fruition.

Finally, a thought also goes out to the administration, the mayor, the aldermen, and the City of Sona officials, for their willingness, their help, and without whom this project would never have been possible.

To everyone, a heartfelt “Thank You”.